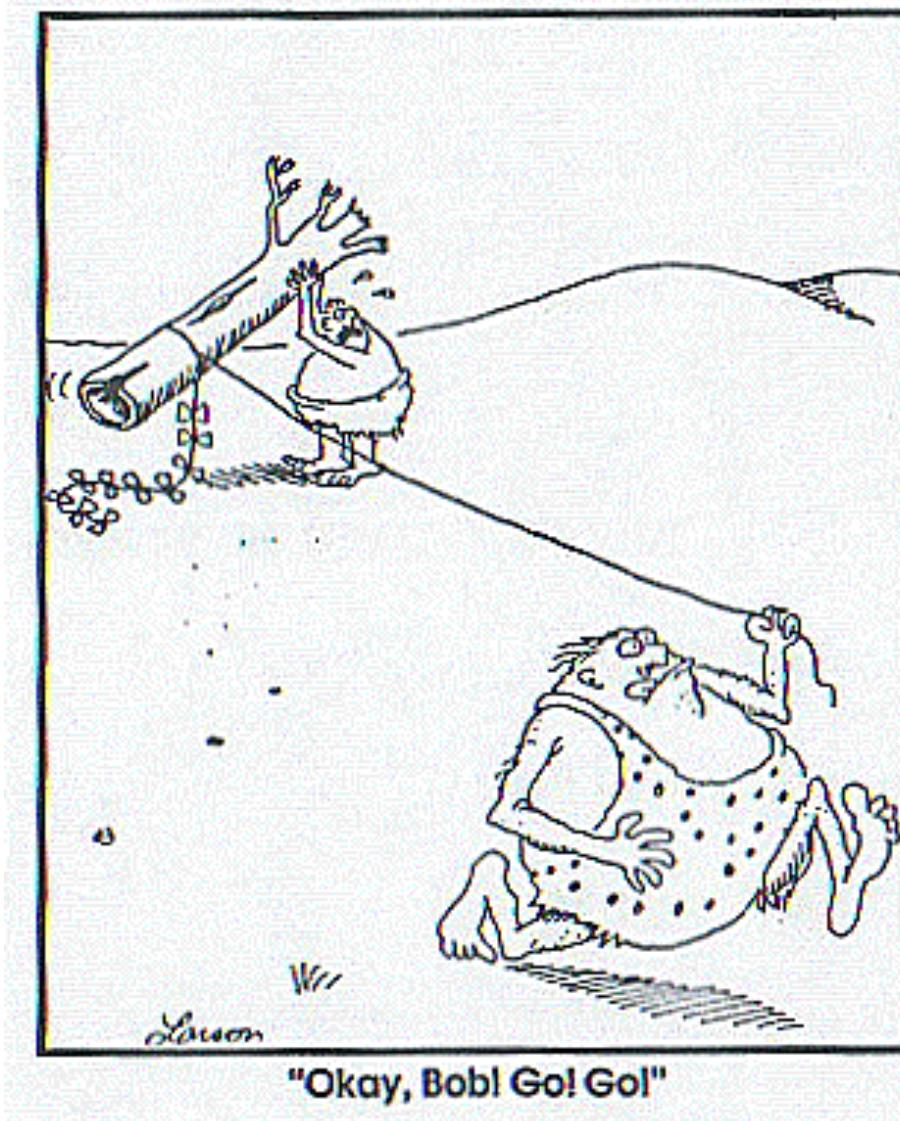


Big things are different from small things



The life time of HEP software

Software is a long-term commitment

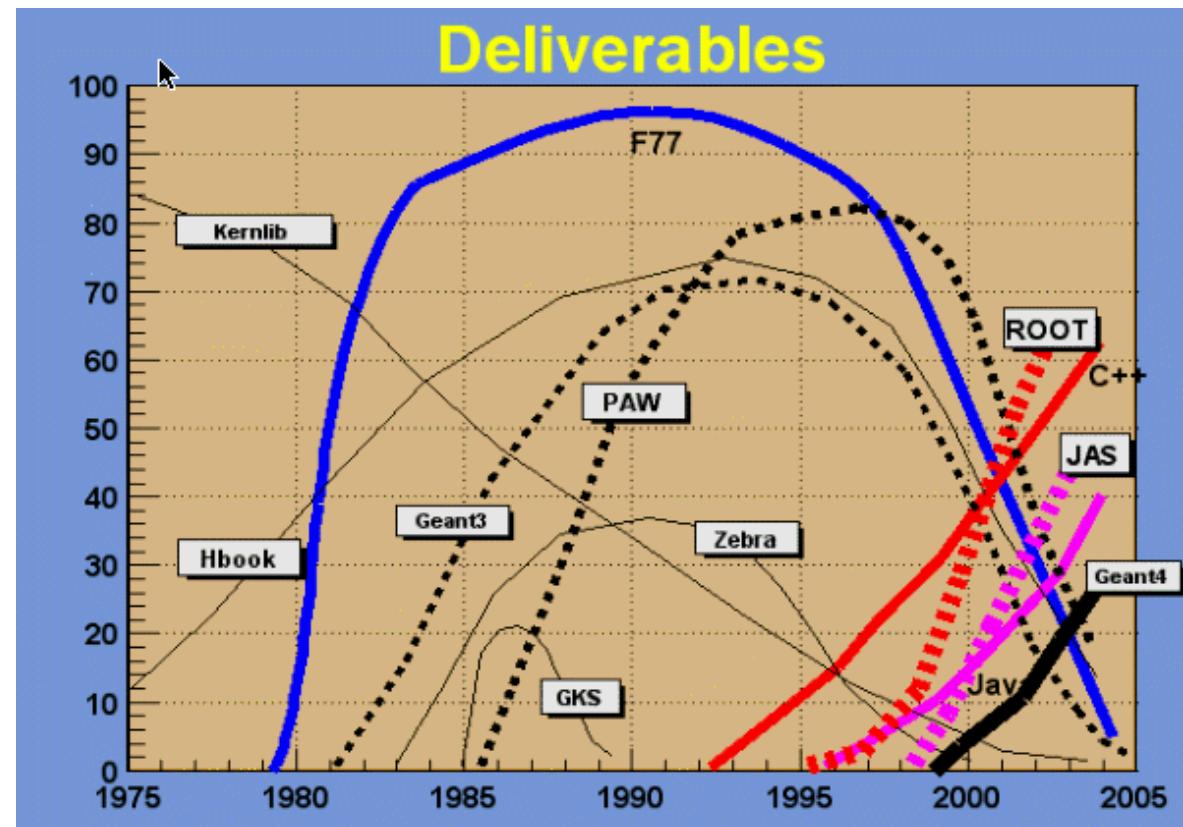
Users like stable and maintained systems

Vote with their feet

It takes time to develop a new system

- Geant3 6+ yrs 3 people 300 KLOCs
- PAW 6+ yrs 5 people 300
- Zebra 4+ yrs 2 people 100
- ROOT 5* yrs 3 people 630

Real work is after that !!



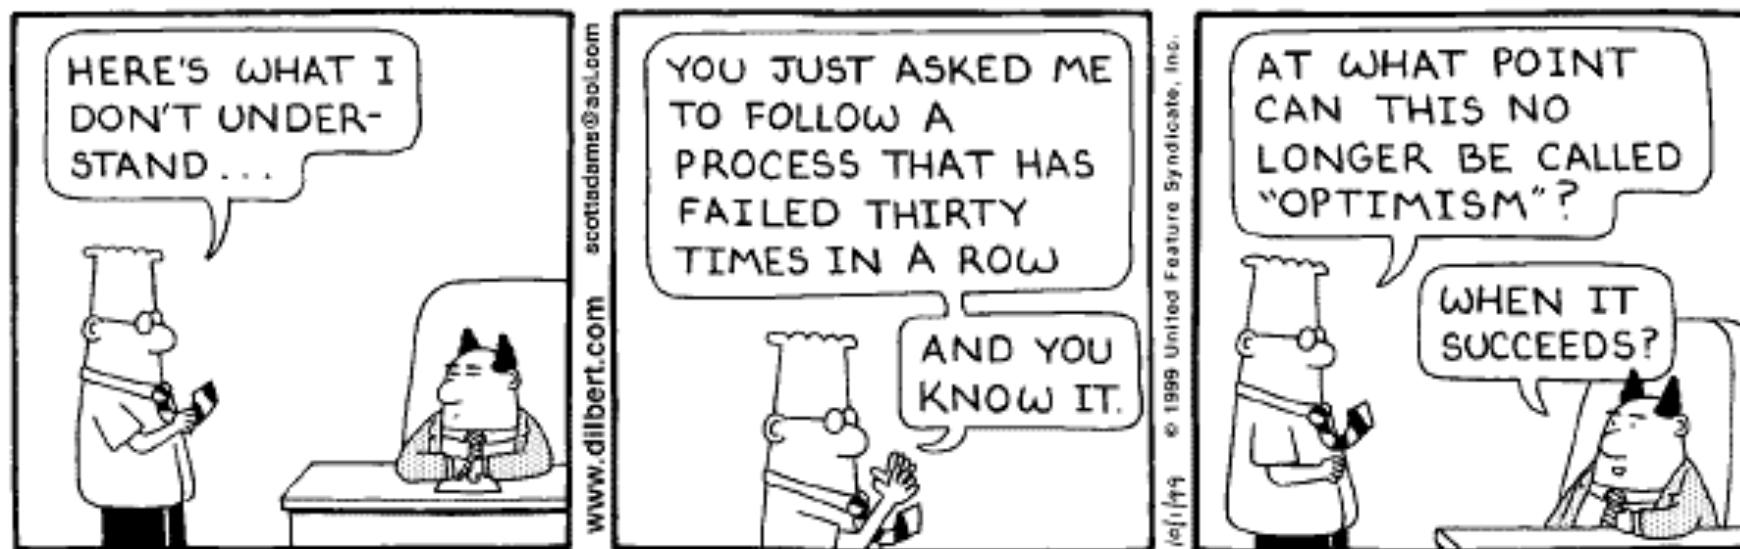
Many releases of the software are needed over its lifetime
to fix bugs, add new features, support new platforms etc

How do we cope?

We try to find a way of working that leads to success

- We create a “process” for building systems
- We devise methods of communicating and record keeping: “models”
- We use the best tools & methods we can lay our hands on

And we engage in denial:



Can't technology save us?

We've built a series of ever-larger tools to handle large code projects:

- CVS, SVN, Git for controlling and versioning code
- Tools for building “releases” of systems
- Tools for “configuration management”

More

But we struggle against three forces:

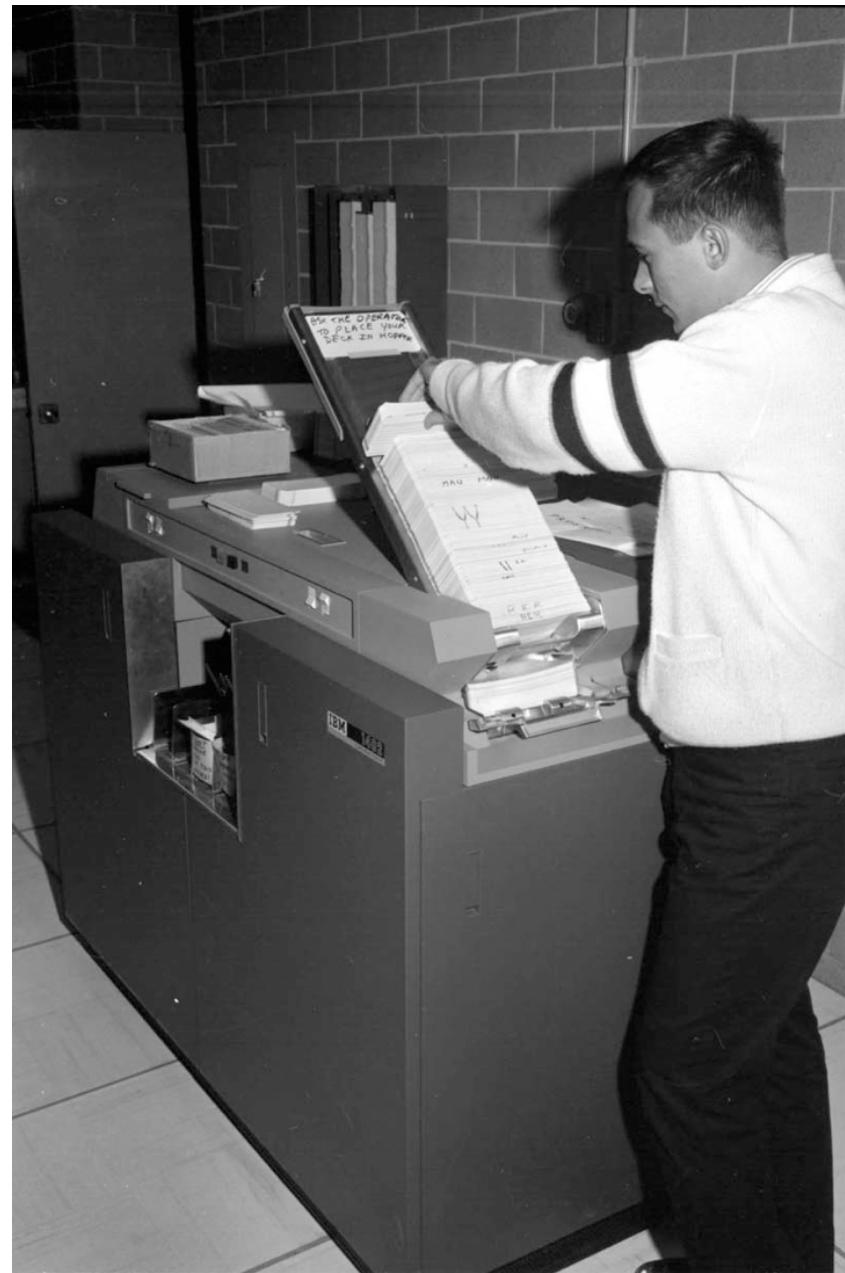
- We're always building bigger & more difficult systems
- We're always building bigger & more difficult collaborations
- And we're the same old people

Net effect: We're always pushing the boundary of what we can do

Stupidity got us into this mess; why can't it get us out? - Will Rogers

How we got here:

First, you just wrote a big program



How we got here:

First, you just wrote a big program

But soon it was so big you wanted help



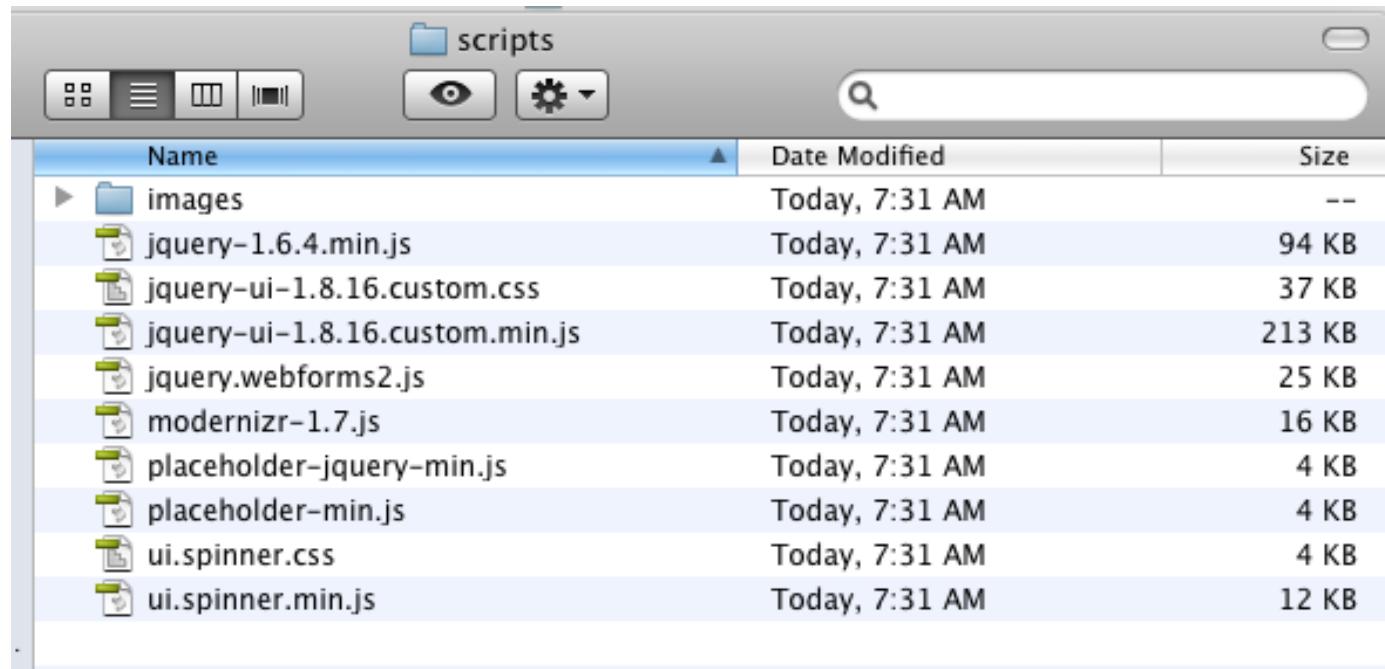
Bob Jacobsen, UC Berkeley

How we got here:

First, you just wrote a big program

But soon it was so big you wanted help

So you broke it into pieces/files/modules



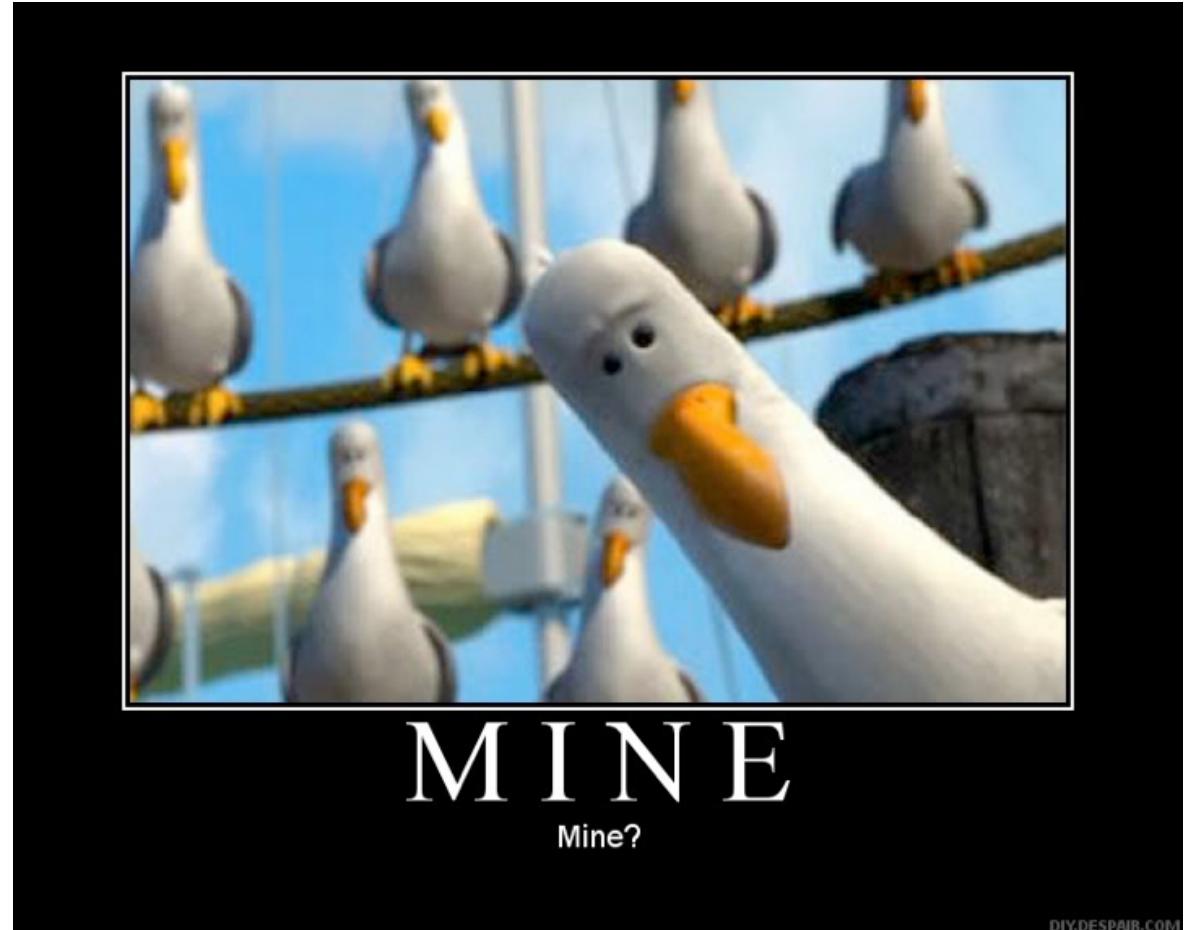
How we got here:

First, you just wrote a big program

But soon it was so big you wanted help

So you broke it into pieces/files/modules

But how do you share work on those?



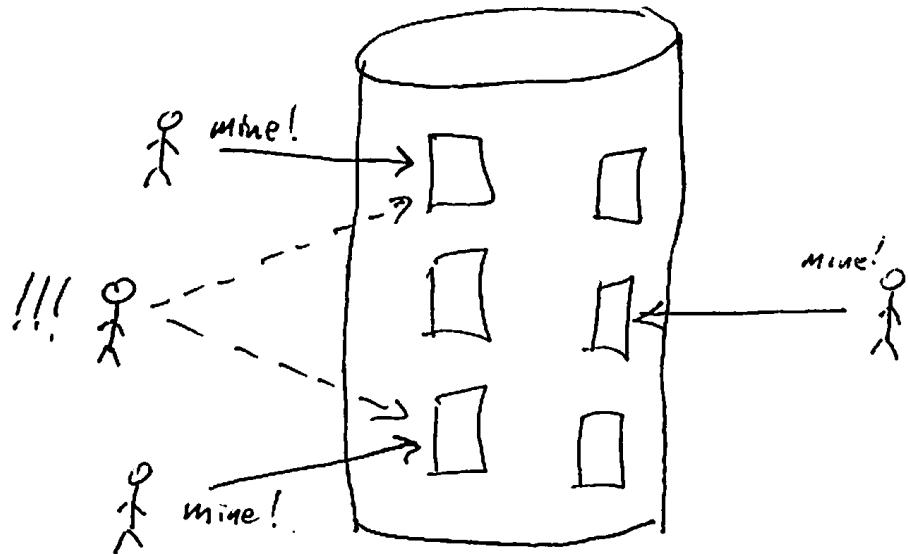
How we got here:

First, you just wrote a big program

But soon it was so big you wanted help

So you broke it into pieces/files/modules

But how do you share work on those?

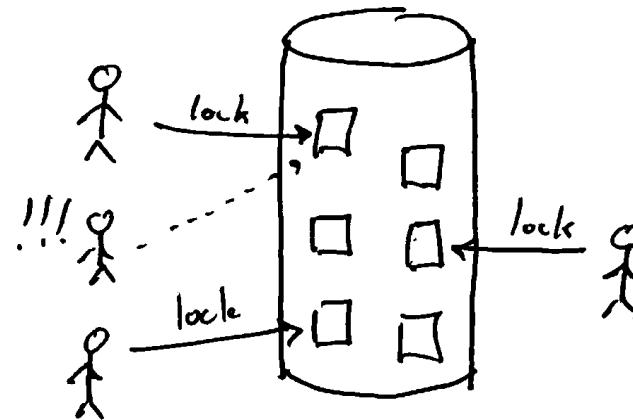


Revision Control System (RCS)

Maintains a repository of text files

- Allows users to check-out, edit, check-in changed text
- Old code remains available
 - Each checked-in change defines a new revision
 - You can retrieve, ask for differences with any of them
- Revisions can be tagged for easy reference

Anybody can get a specific set of source code file versions

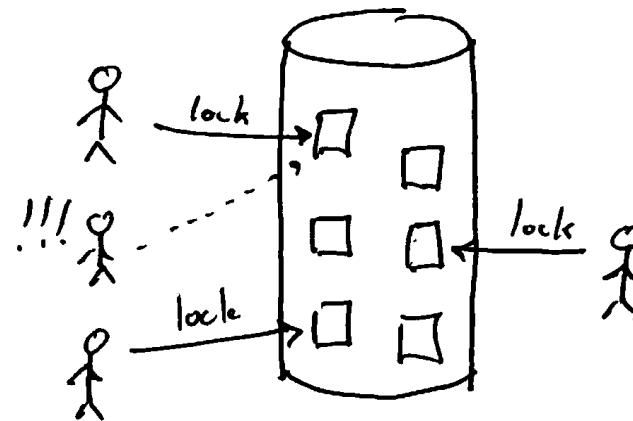


Revision Control System (RCS)

Maintains a repository of text files

- Allows users to check-out, edit, check-in changed text
- Old code remains available
 - Each checked-in change defines a new revision
 - You can retrieve, ask for differences with any of them
- Revisions can be tagged for easy reference

Anybody can get a specific set of source code file versions



But only one person working on a file at a time!

Problem: This serializes development

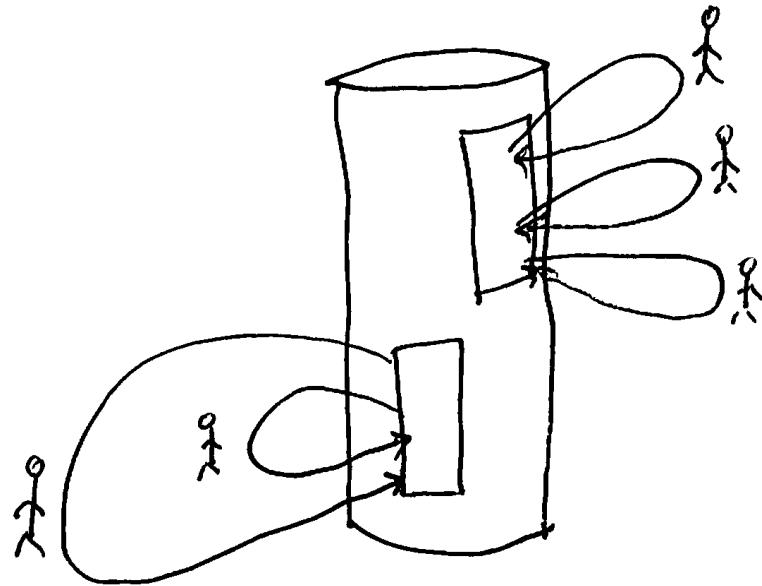
Workarounds, but with problems of their own

Concurrent Versions System (CVS)

As systems & collaborations grow, efficiency goes down

More

“Version” idea: Track changes from one version to next



Big advantage: checkout is not exclusive

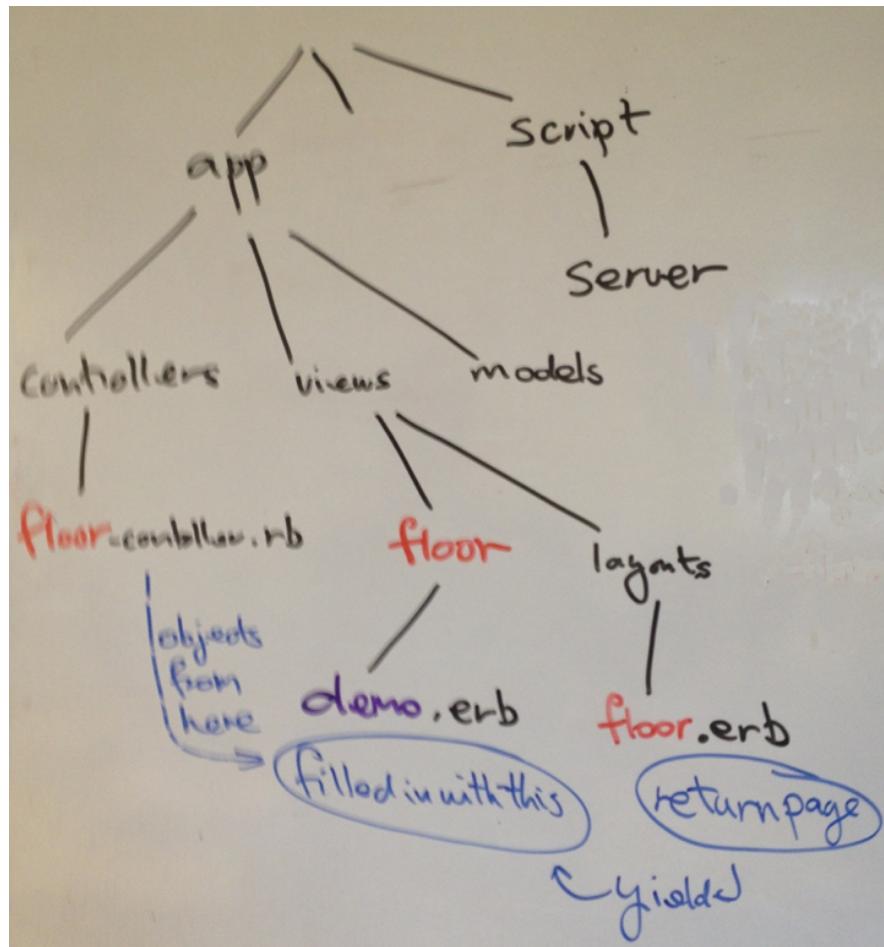
- More than one developer can have the same file checked out
- Developers can control their own use of the code for read, write
- Changes can come from multiple sources
- Tool handles (most) of the conflict resolution

And systems still grow

You broke the code into pieces/files/modules

And things got more and more complicated

You needed an organization above the level of the file



Directory Tree

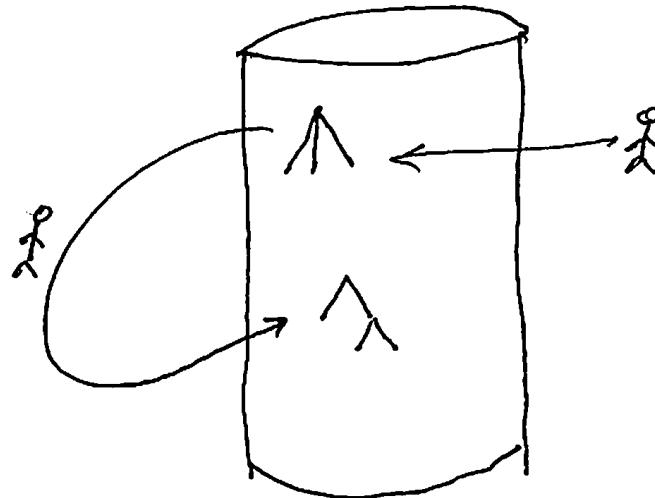
```

.
|-- code
|   |-- perl
|   |   |-- clean_fasta.pl
|   |   `-- run_glam2.pl
|   |-- R
|   |   |-- gse21350.R
|   |   |-- gse7275.R
|   |   `-- gse9589.R
|   |-- ruby
|   |   '-- seq2svm.rb
|   '-- sh
|       '-- ens_uniq.sh
|-- data
  
```

Subversion (svn)

So you broke it into pieces/files/modules
And things got more and more complicated
You needed an organization above the level of the file
Want to be able to collaborate on that:

More



Subversion (svn) brings tools for doing that

Why isn't that enough?

CVS, SVN lets me “check out” complete source code. Then just compile!

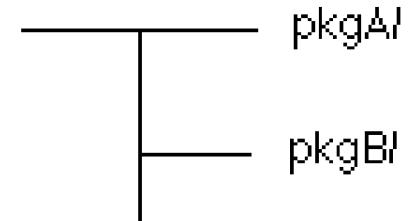
- Works great for small projects
- But builds run into several levels of scaling problems

1) Want to attach to external code

- We don't write everything (though tempted)
- Sometimes don't get source for external code
- Need some way to connect to specific external libraries:
Both specific product, and a specific version of that product

2) Want to separate code into multiple parts

- So people/institutions can take responsibility for parts
- But software has cross-connections
- Need structure that works for both

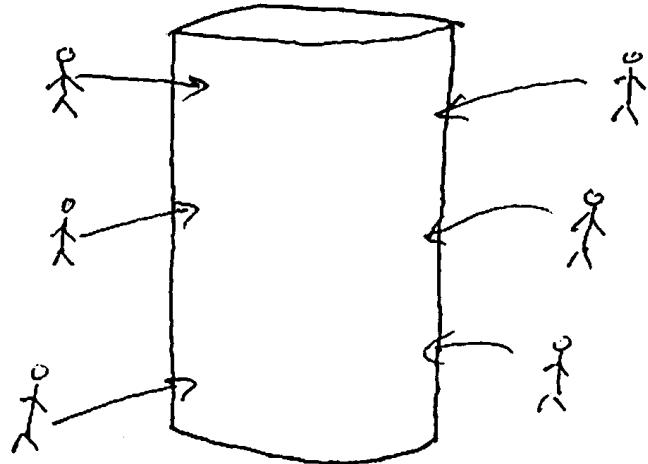


And still need to be able to build the code

Scaling is still an issue

Everybody is sharing a single repository

Every commit is immediately visible to everybody else



Development stands on shifting sand

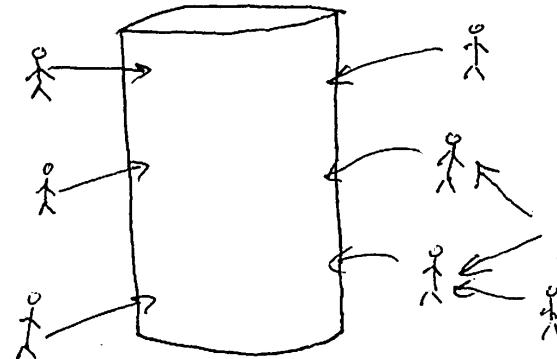
Detailed records, but little understanding

Workarounds!

Package Coordinators

Tags and Branches

External record keeping tools



Scaling: Handling complicated builds

[More](#)

Multiple “packages” require cross connects while compiling

- Typing the compile command gets boring fast

```
g++ -c -I"/afs/cern.ch/user/s/scherzer/public/1001/InstallArea/include/PixelDigitization"
-I"/afs/cern.ch/user/s/scherzer/public/1001/InstallArea/include/SiDigitization"
-I"/afs/cern.ch/atlas/software/dist/10.0.1/InstallArea/include/InDetSimEvent"
-I"/afs/cern.ch/atlas/software/dist/10.0.1/InstallArea/include/HitManagement"
-I"/afs/cern.ch/atlas/software/dist/10.0.1/InstallArea/include/TestTools"
-I"/afs/cern.ch/atlas/software/dist/10.0.1/InstallArea/include/TestPolicy"
-I"/afs/cern.ch/atlas/offline/external/Gaudi/0.14.6.14-pool201/GaudiKernel/v15r7p4"
-I"/afs/cern.ch/sw/lcg/external/clhep/1.8.2.1-atlas/slc3_ia32_gcc323/include"
-I"/afs/cern.ch/sw/lcg/external/Boost/1.31.0/slc3_ia32_gcc323/include/boost-1_31"
-I"/afs/cern.ch/sw/lcg/external/cernlib/2003/slc3_ia32_gcc323/include" -O2 -pthread
-D_GNU_SOURCE -pthread -pipe -ansi -pedantic -W -Wall -Wwrite-strings -Woverloaded-virtual
-Wno-long-long -fPIC -march=pentium -mcpu=pentium -pedantic-errors -ftemplate-depth-25
-ftemplate-depth-99 -DHAVE_ITERATOR -DHAVE_NEW_IOSTREAMS -D_GNU_SOURCE
-o PixelDigitization.o -DEFL_DEBUG=0 -DHAVE_PRETTY_FUNCTION -DHAVE_LONG_LONG
-DHAVE_BOOL -DHAVE_EXPLICIT -DHAVE_MUTABLE -DHAVE_SIGNED -DHAVE_TYPENAME
-DHAVE_NEW_STYLE_CASTS -DHAVE_DYNAMIC_CAST -DHAVE_TYPEID
-DHAVE_ANSI_TEMPLATE_INSTANTIATION -DHAVE_CXX_STDC_HEADERS '
-DPACKAGE_VERSION="PixelDigitization-00-05-16" -DNDEBUG -DCLHEP_MAX_MIN_DEFINED
-DCLHEP_ABS_DEFINED -DCLHEP_SQR_DEFINED ..../src/PixelDigitization.cxx
```

Build tools: “make”, “Ant”, etc

- Manually create a “makefile” that forwards include options to the compiler

```
g++ -IpkgA -IpkgB
```

- Lets you adapt to various internal structures

```
g++ -IpkgA -IpkgB/include -IpkgC/headers
```

- Also lets you add other options to control localization, debugging, etc

Size keeps getting in the way

Small experiment (offline production code only):

- 430 packages
- 17,000 files
- 7 million lines of source

Some of these are large “for historical reasons”

But that’s true of just about any project

Repository checkout: 13 minutes

Build from scratch: 6 hours

Spread across multiple production machines; never did complete on laptop

“gmake” with one change: about 4-12 minutes to think about dependencies

And I don’t even want to think about the size of a monolithic Makefile

And everybody will need multiple copies...

Old ones, new ones, ...

“But I just want to run the program!”

Issue arises at large & small level

At the level of developers, need way to manage this

- Both tools and procedures

We'll be discussing & exercising typical tools; many exist!

Individual collaborations have their own ways of sharing info

At the collaboration level, need procedures to ensure it works

- “Nightly builds”

Now common in HEP - Give early feedback on consistency problems

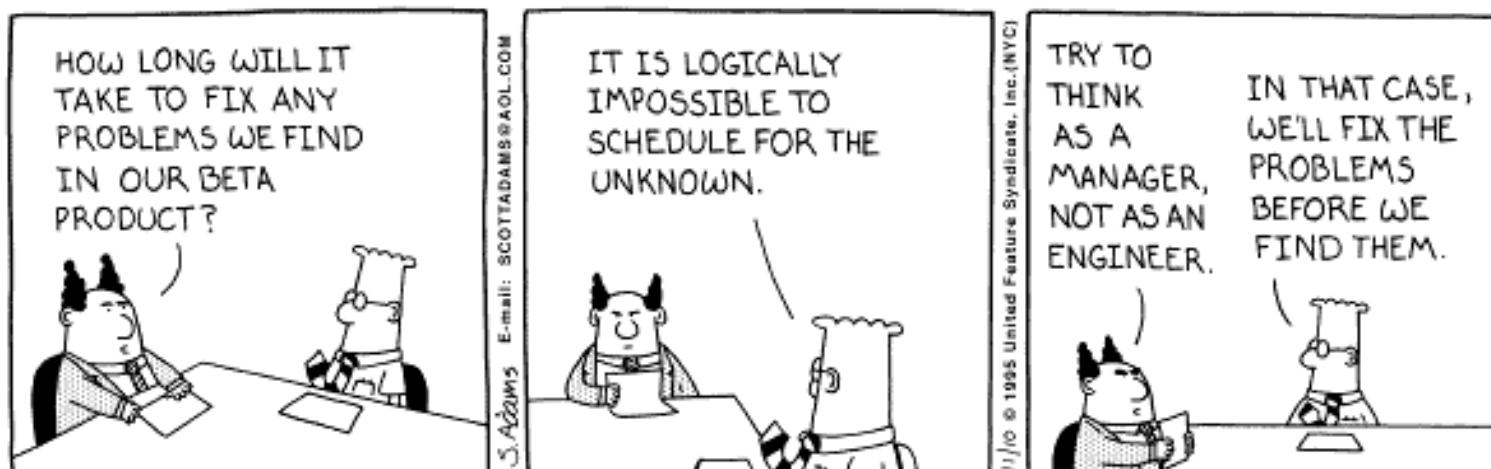
Many in industry moving toward “continuous integration”

More

- Not a complete solution by itself

Only works when people actually integrate early and often

- Reduces problems, but integration is still a lot of work



When Boeing wanted to design the 747, they had two choices:

1. Hire “SuperEngineer”, who could do it alone
2. Hire 7,200 engineers and organize them to cooperate

Which did they choose?

Why?

What can we learn from this?



Two Approaches: (1) Organize people to match the work

1) You want a “release system”

Use tags in the repository to mark “package versions”

“Package Coordinators” are people with the local knowledge

Build tools that record relations between packages, external requirements:

- pull out proper consistent versions,
- combine make files,
- control the build

Complicated tools that need to know a bunch of stuff

Configuration Management Tool (CMT)

- Based on ‘requirements file’ with custom syntax and contests

Lots of others (SCRAM, ETICS, cloud-based tools)

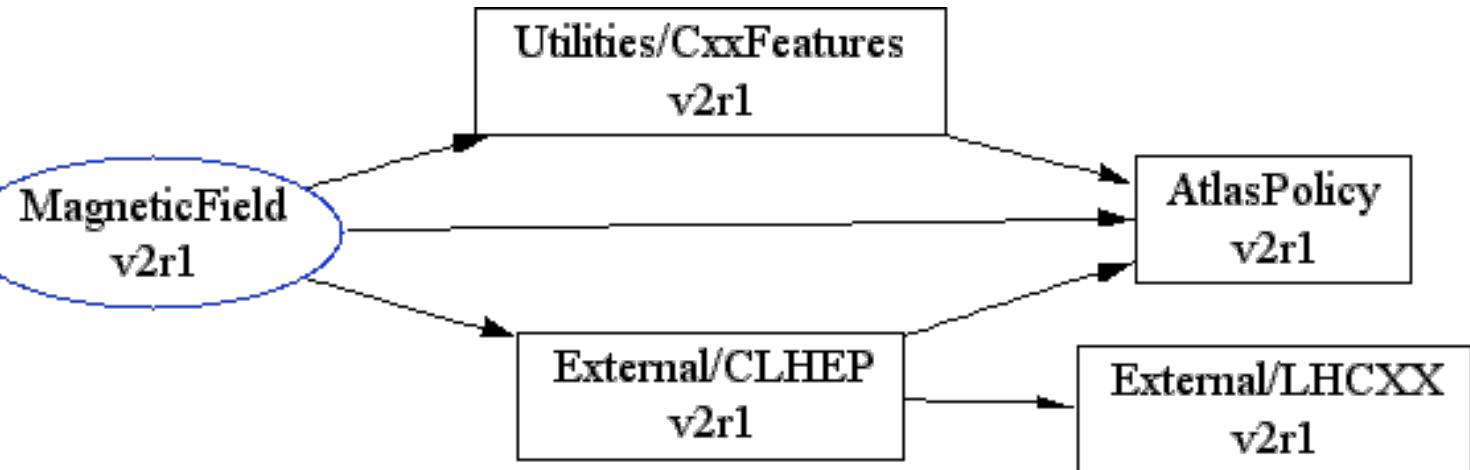
- Showing CMT because easy to see how it works

More

CMT: a simple release and consistency tool

Requirements file provides custom language for expressing our needs

More



```

package MagneticField

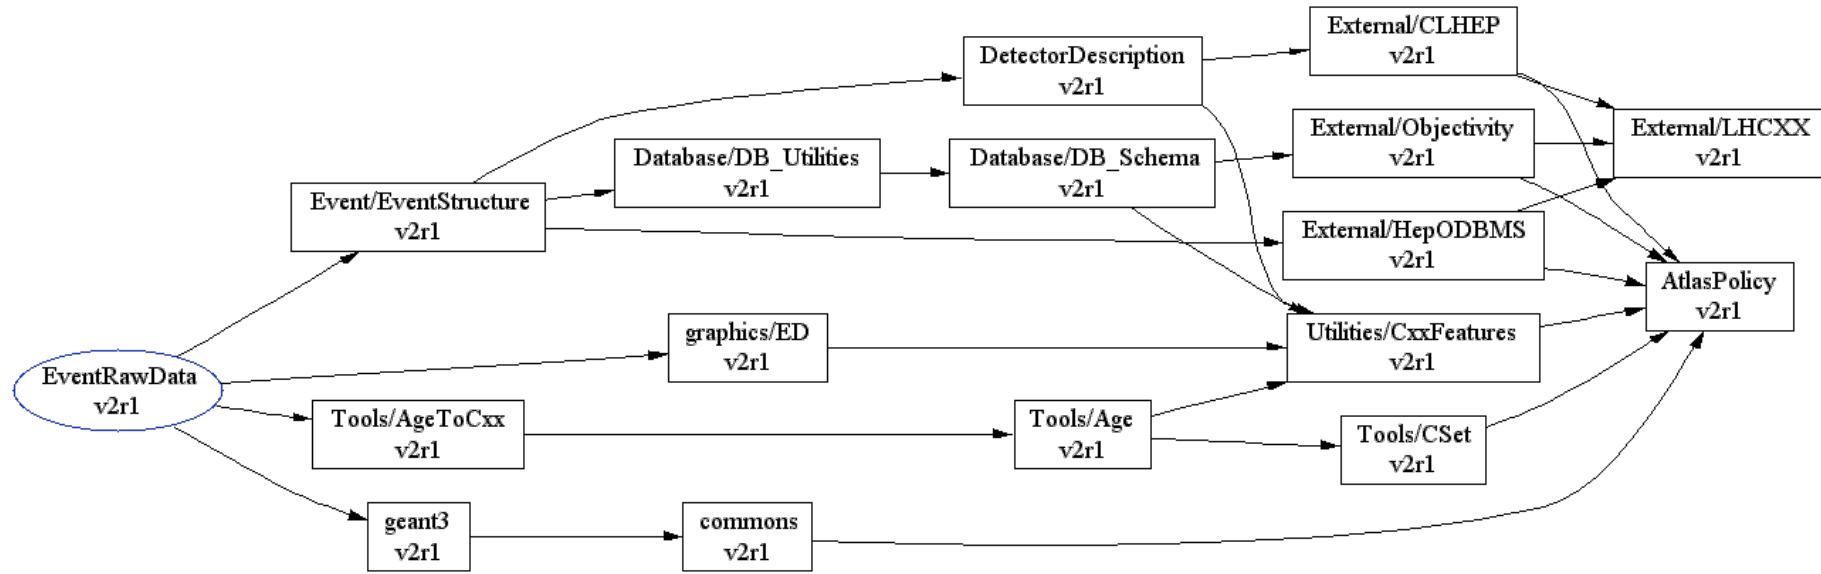
author Laurent Chevalier <laurent@hep.saclay.cea.fr>
author Marc Virchaux <virchau@hep.saclay.cea.fr>

use AtlasPolicy v2r1
use CxxFeatures v2r1 Utilities
use CLHEP v2r1 External

include_dirs $(MAGNETICFIELDROOT) /MagneticField
branches MagneticField doc src test
...
  
```

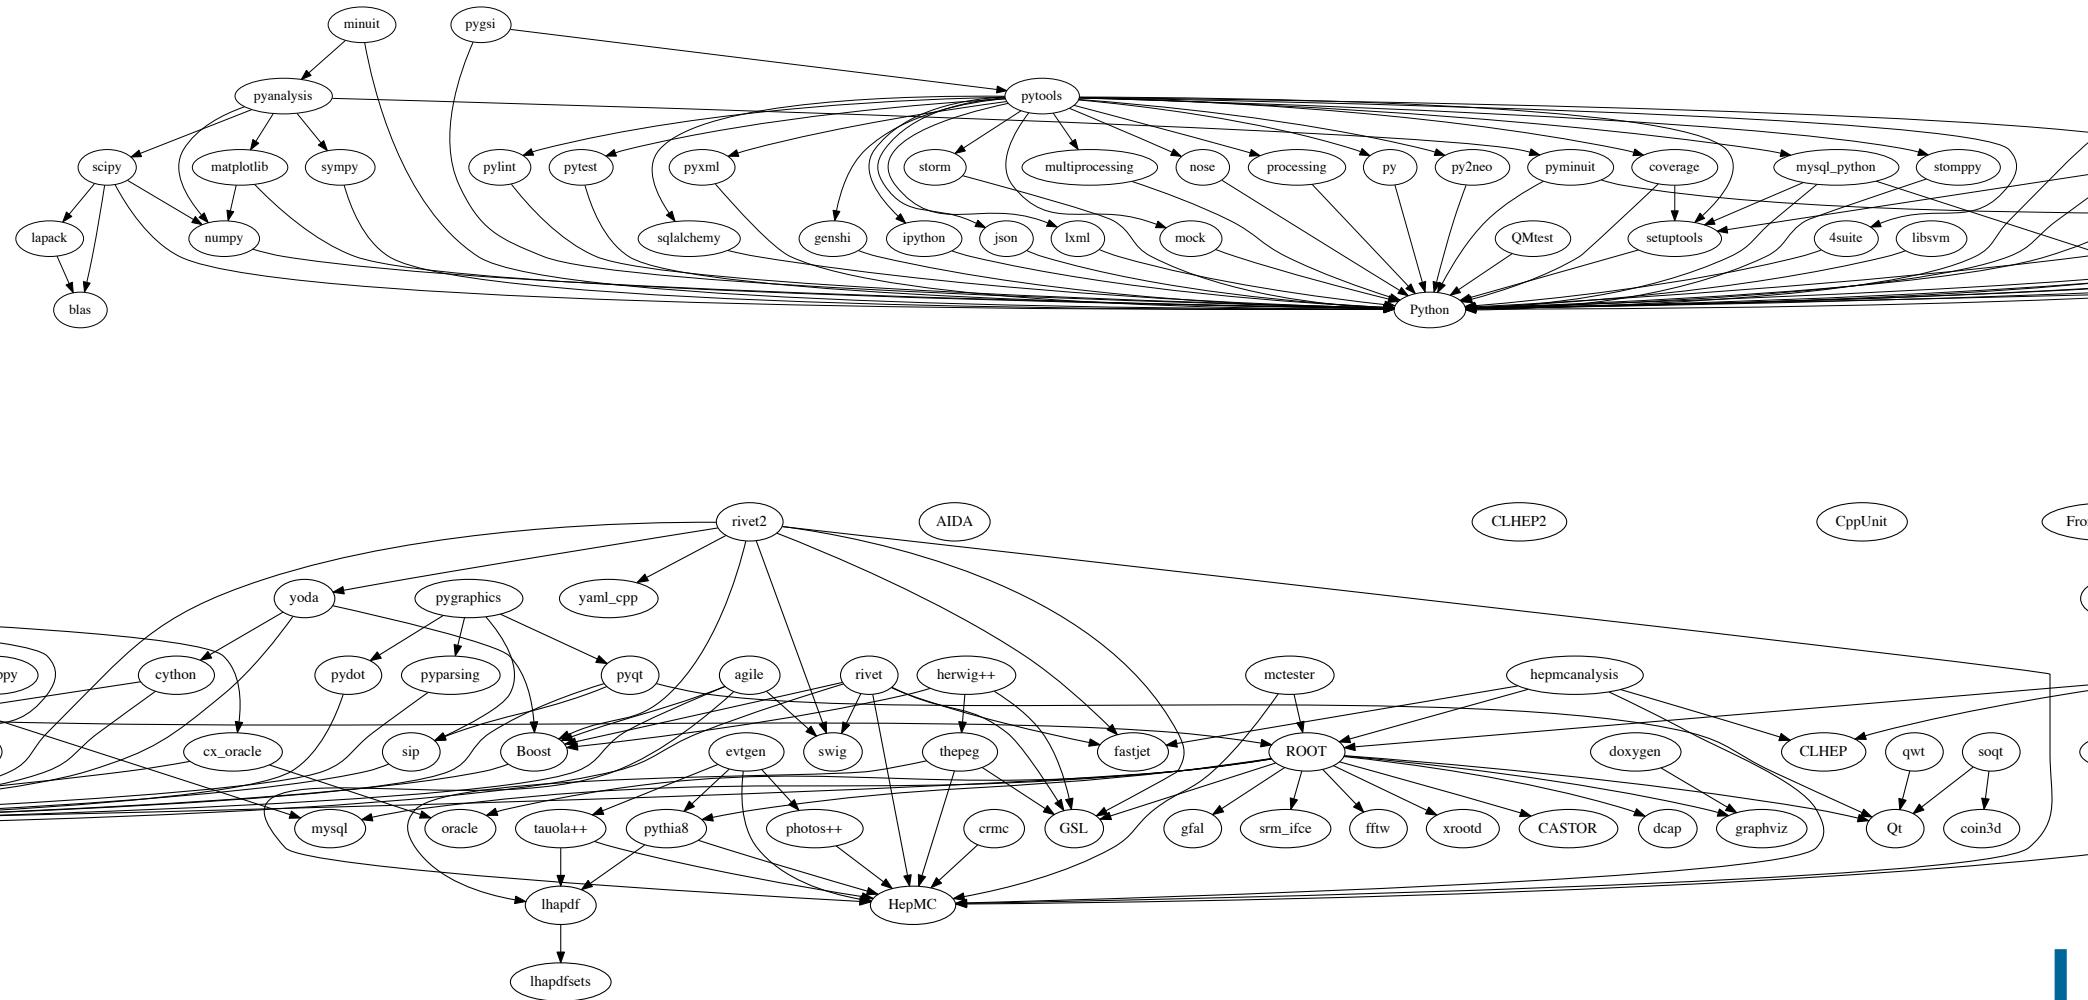
**Example from C.
Arnault (LAL
and Atlas)**

“Consistency” scales poorly



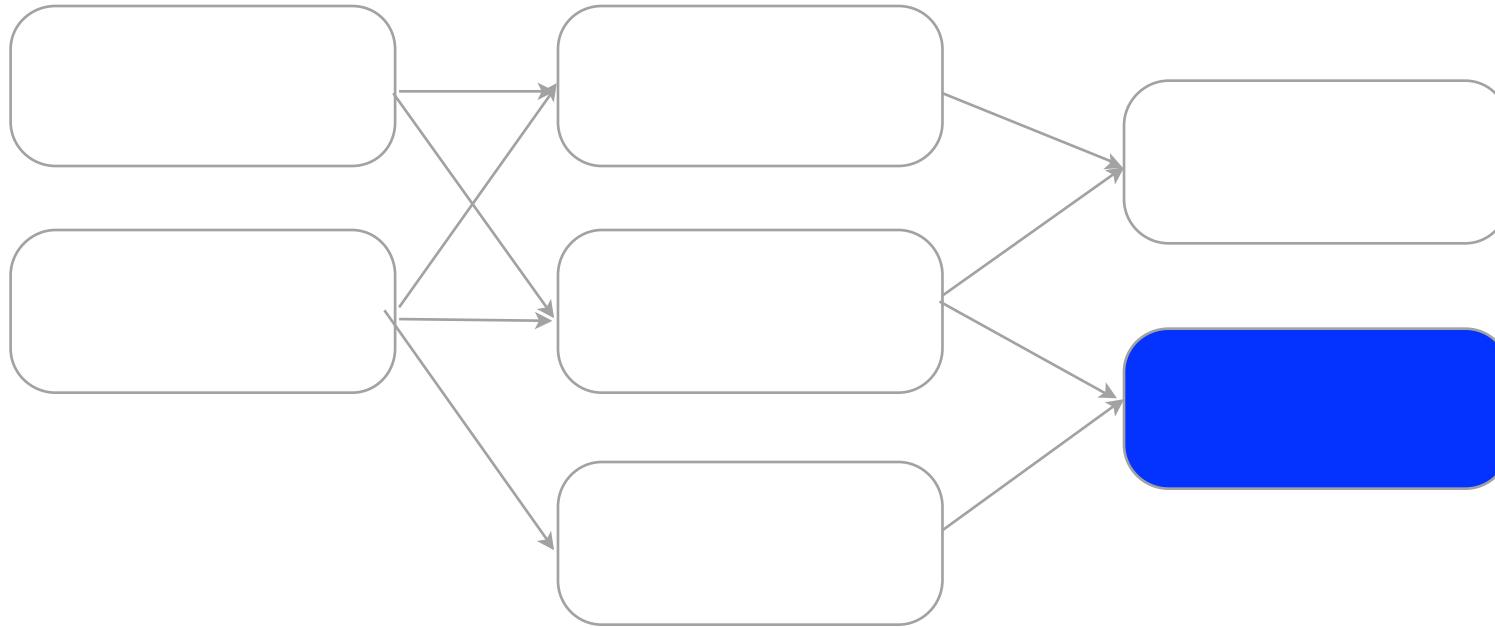
Software strongly depends on other software

- Usually managed at the package level
 - (This can result in lots of packages, as you subdivide over and over)
- Expresses how changes in one piece can drive changes in another

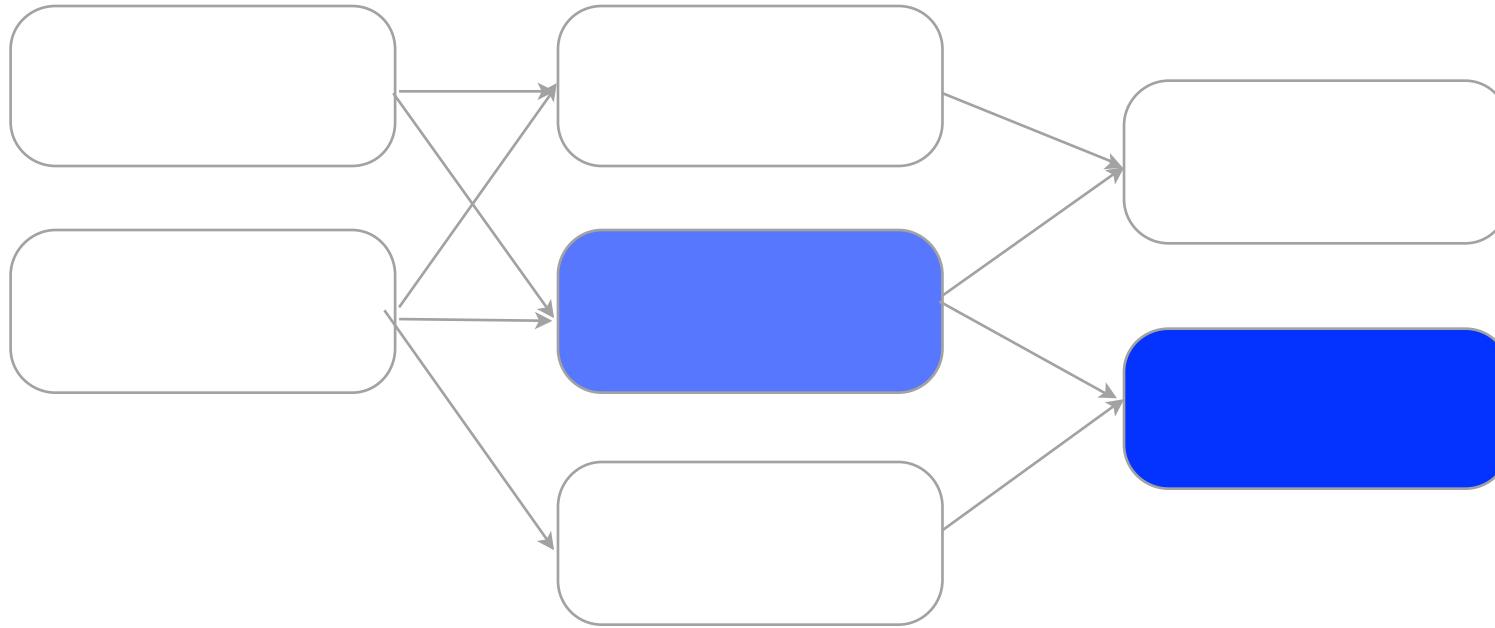


More

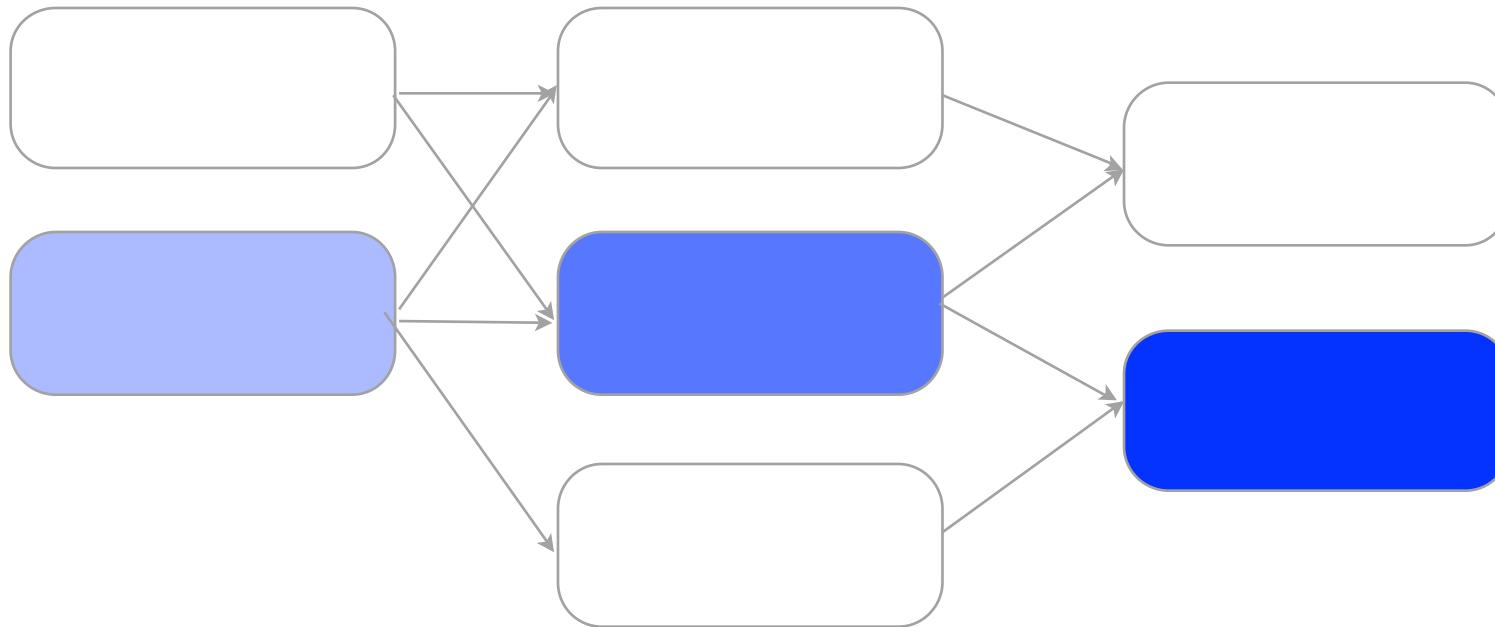
How change propagates through dependencies



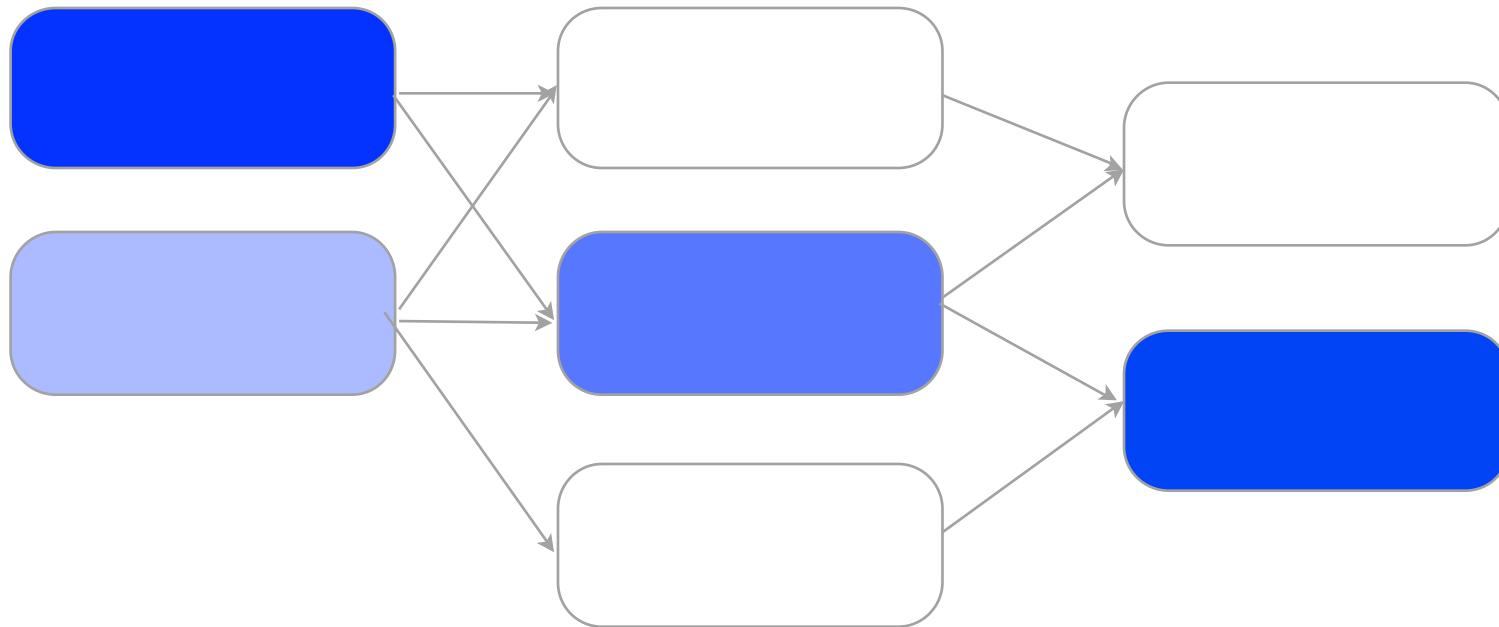
How change propagates through dependencies



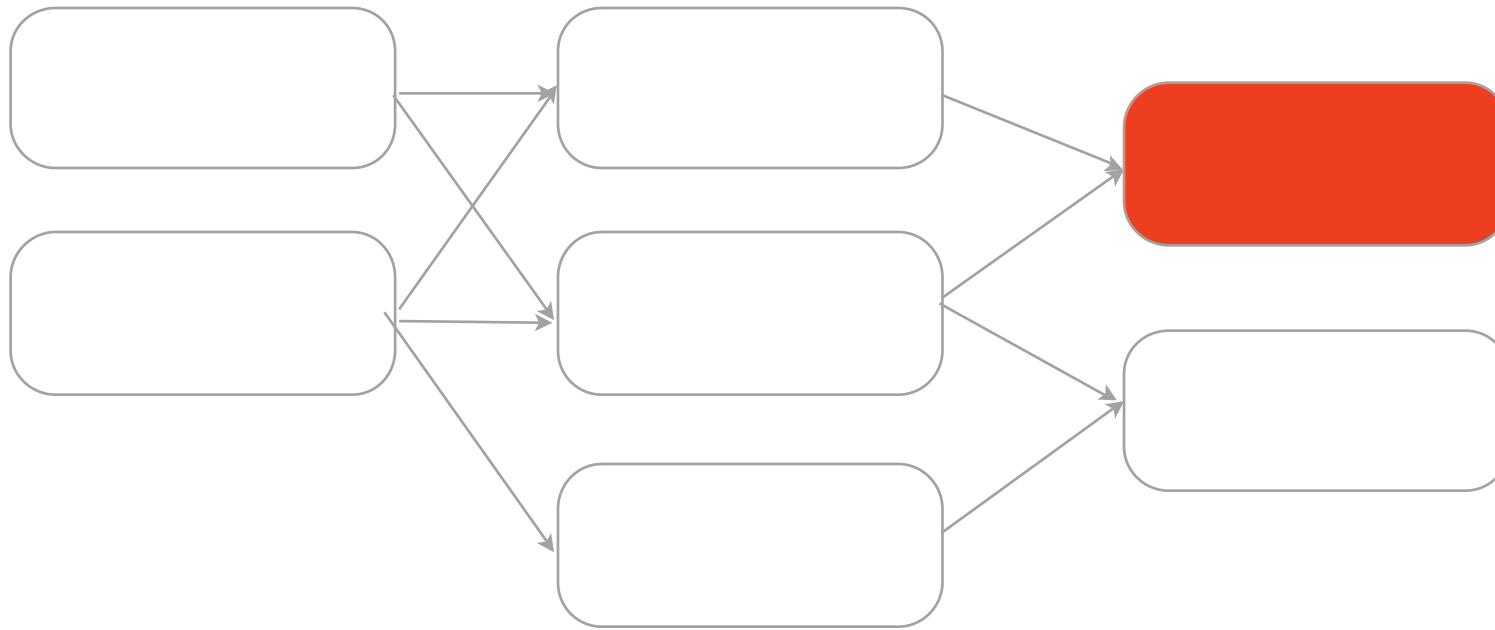
How change propagates through dependencies

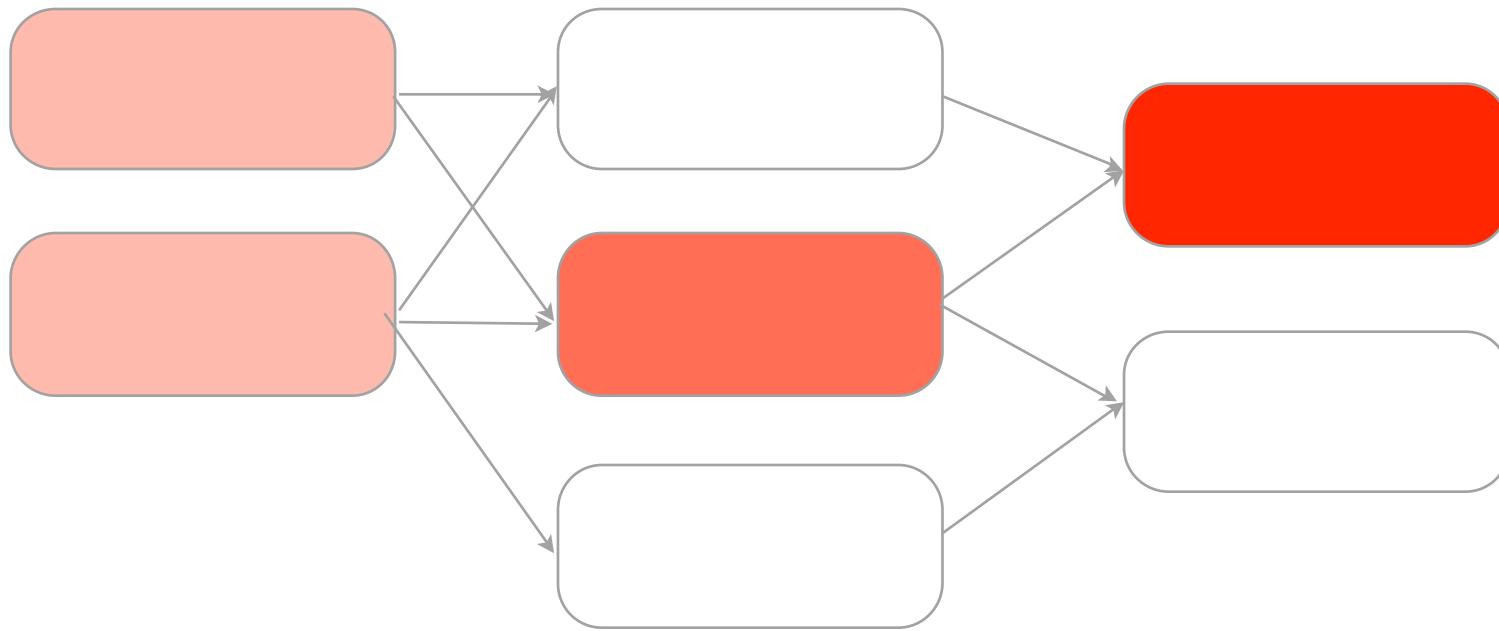


Changes don't always stay small

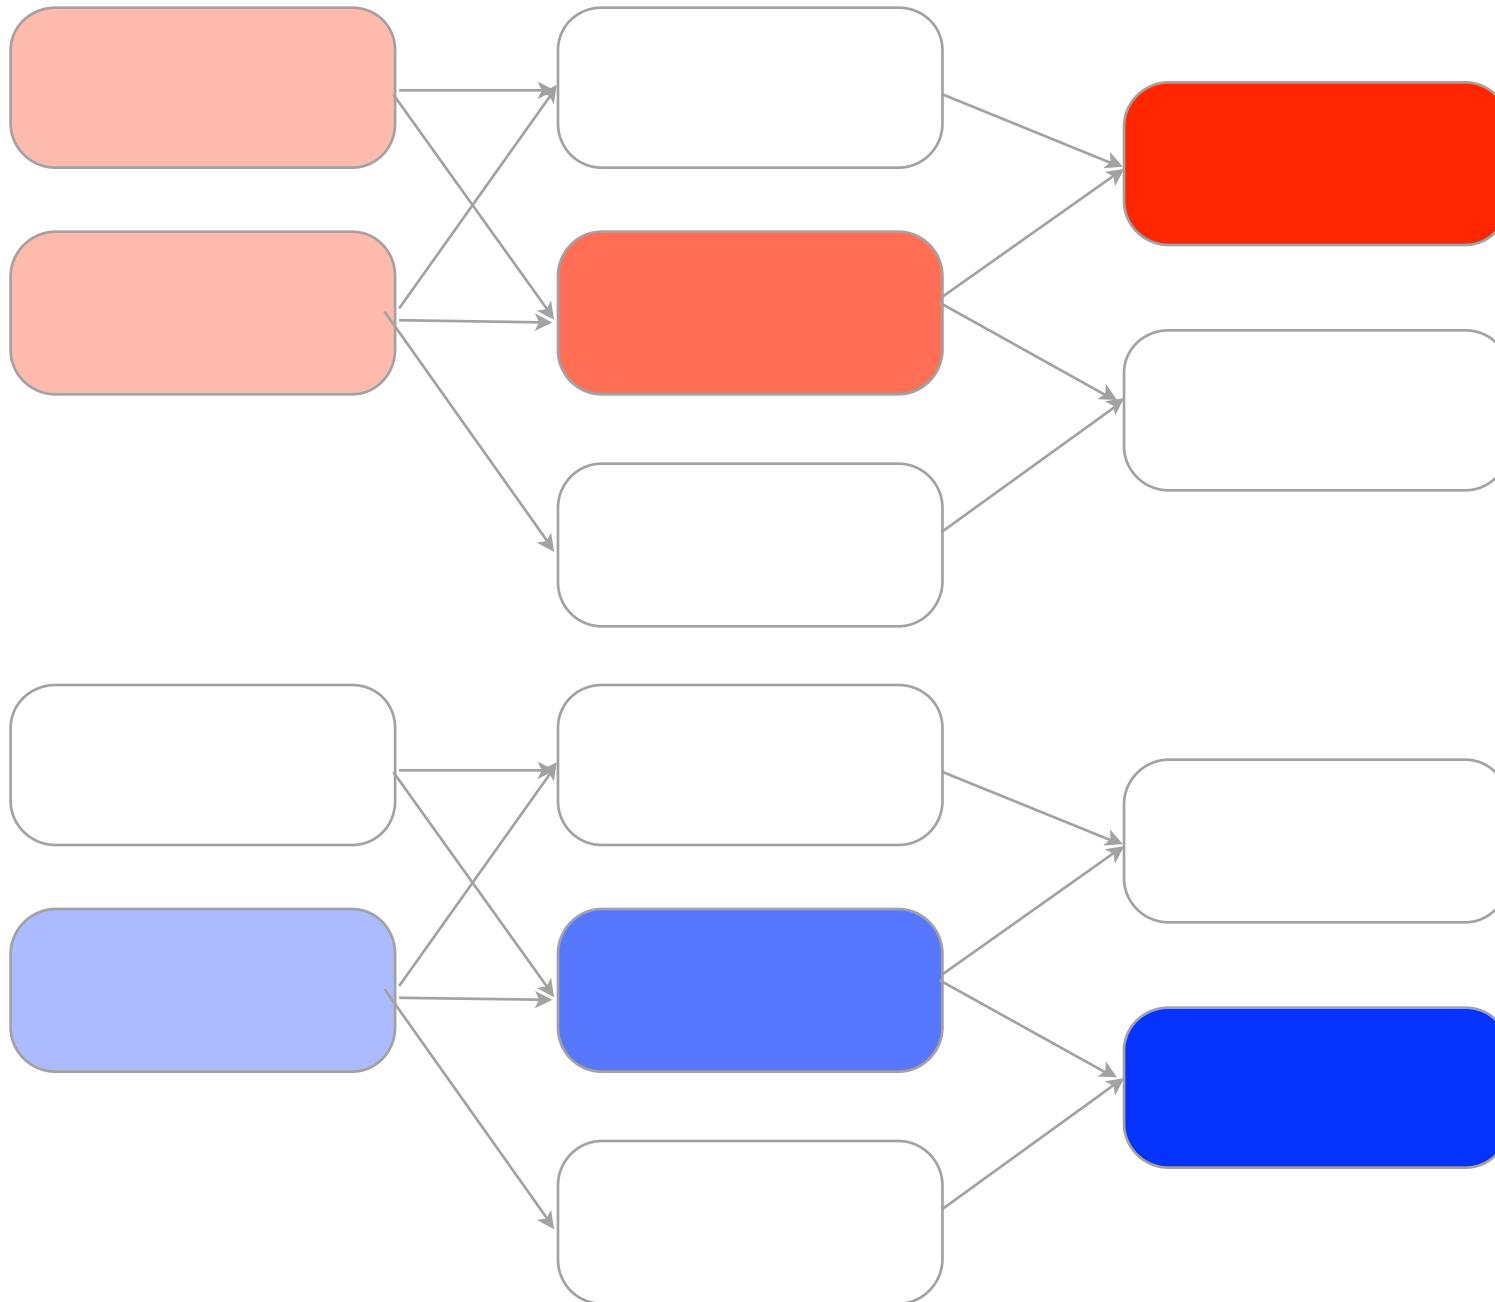


Another change:

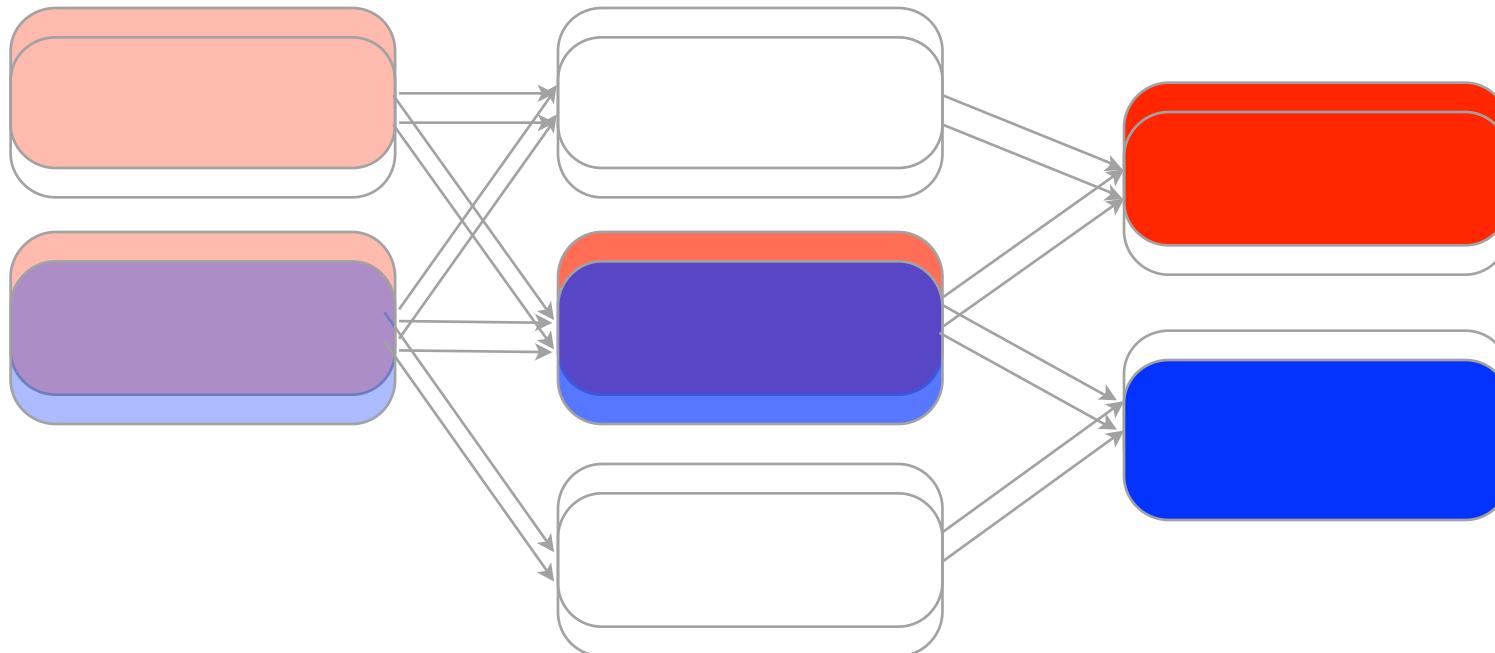




Change management requires people



Change management requires people



2nd approach: People handle consistency, machines build

The repository is where the code should be consistent

Create tools that let you do that!

Allow developers to work on their own content until it's right

- Not every change should go to everybody

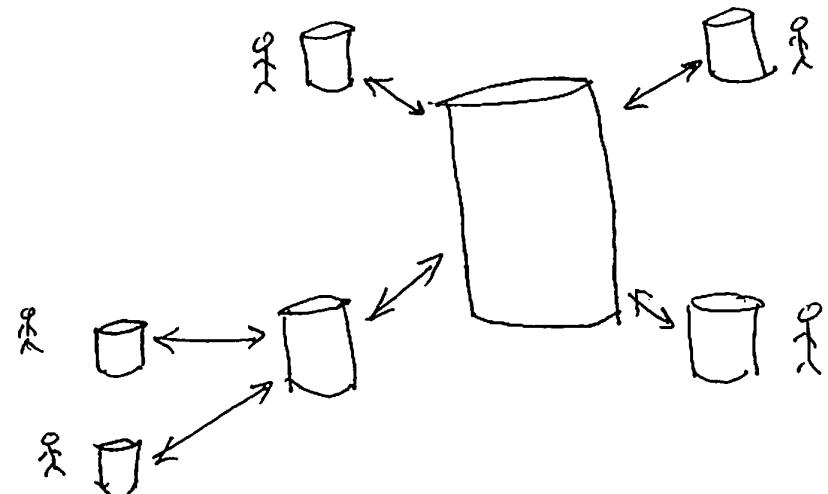
Allow developers to collaborate in small groups

- Put together a sub-system

It's not the versions that are interesting, it's the updates!

- Development becomes a story instead of a series of snapshots
- “Here's our complete contribution”

Enter Git (not an acronym)



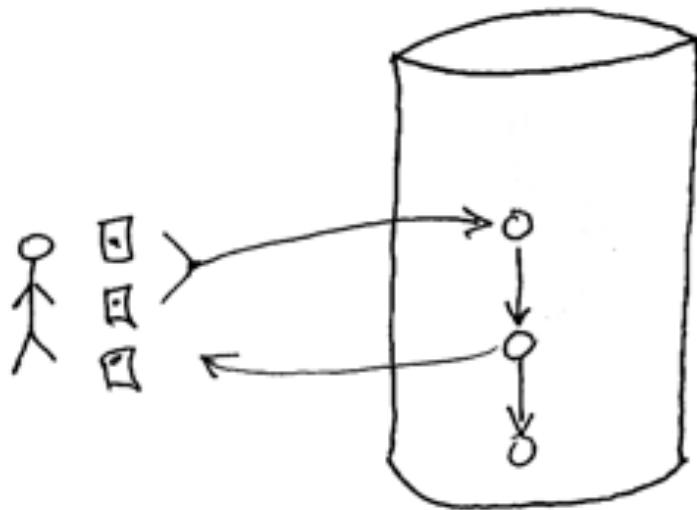
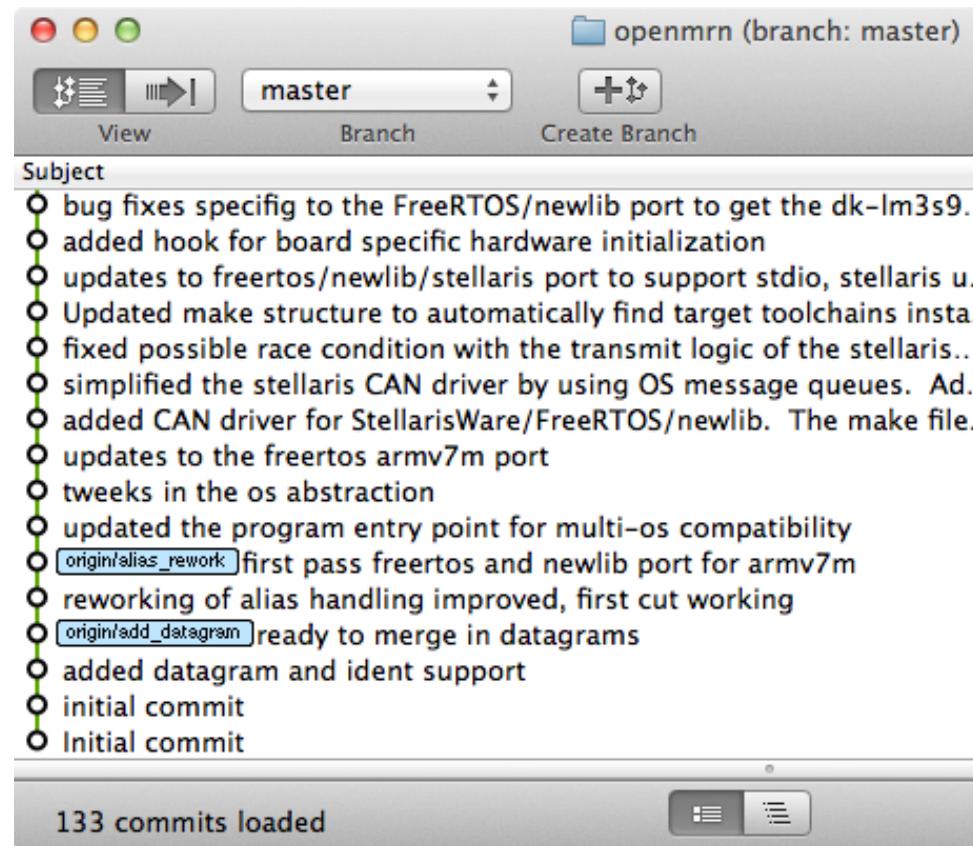
More

At first, Git looks like SVN

You bring out a copy, work on it, and commit

More

Git repository contains all that history

openmrn (branch: master)

master Branch Create Branch

Subject

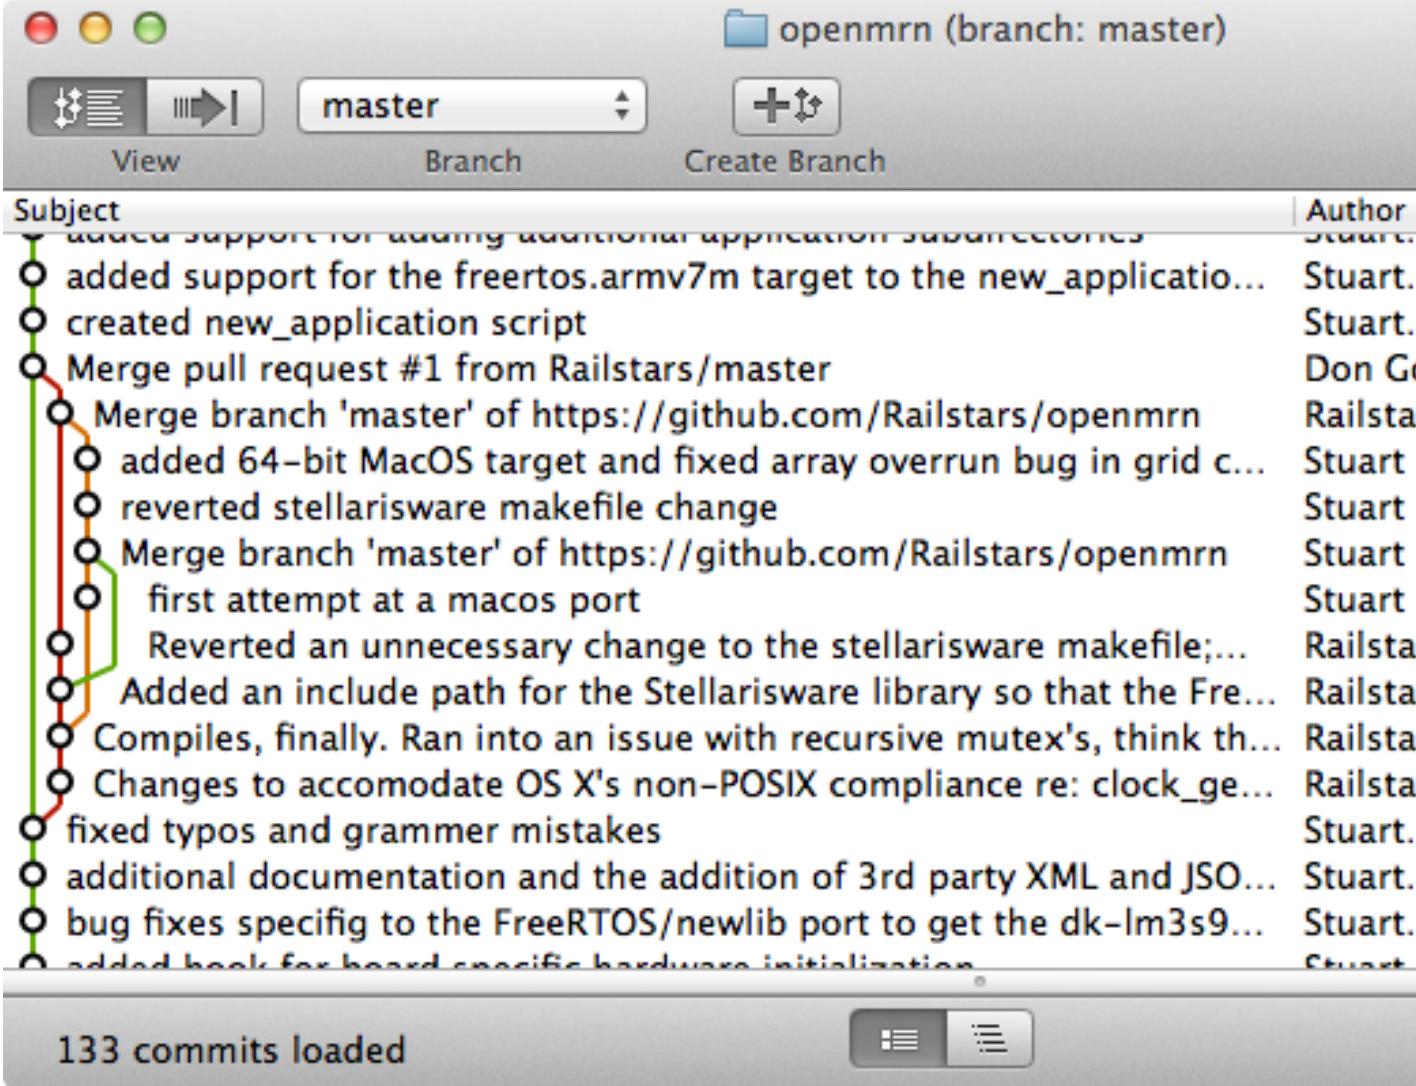
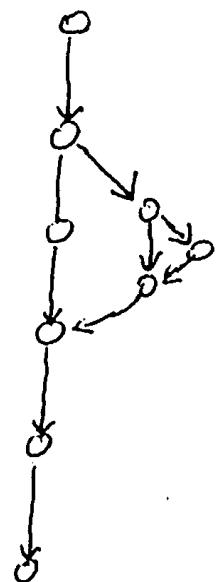
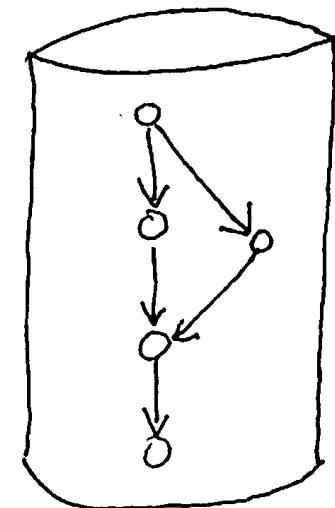
- bug fixes specific to the FreeRTOS/newlib port to get the dk-lm3s9...
- added hook for board specific hardware initialization
- updates to freertos/newlib/stellaris port to support stdio, stellaris u...
- Updated make structure to automatically find target toolchains insta...
- fixed possible race condition with the transmit logic of the stellaris...
- simplified the stellaris CAN driver by using OS message queues. Ad...
- added CAN driver for StellarisWare/FreeRTOS/newlib. The make file...
- updates to the freertos armv7m port
- tweaks in the os abstraction
- updated the program entry point for multi-os compatibility
- origin/alias_rework first pass freertos and newlib port for armv7m
- reworking of alias handling improved, first cut working
- origin/add_datagram ready to merge in datagrams
- added datagram and ident support
- initial commit
- Initial commit

133 commits loaded

“Scratchpad” idea lets you control what you commit: Shaping the story

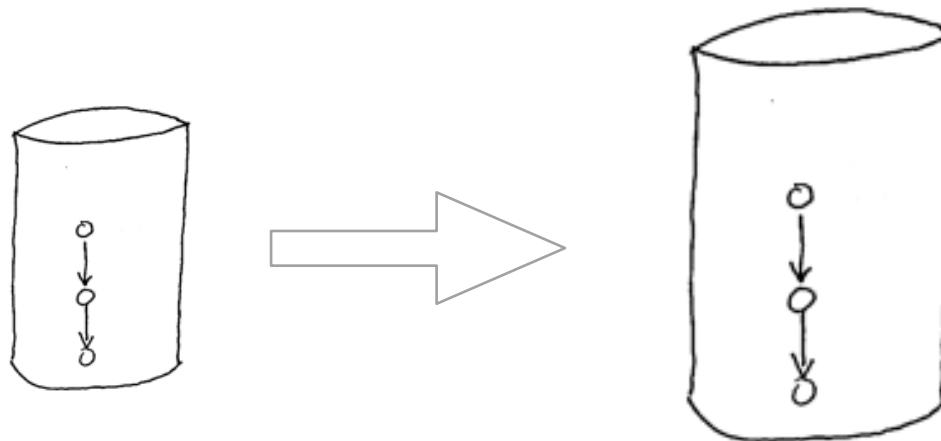
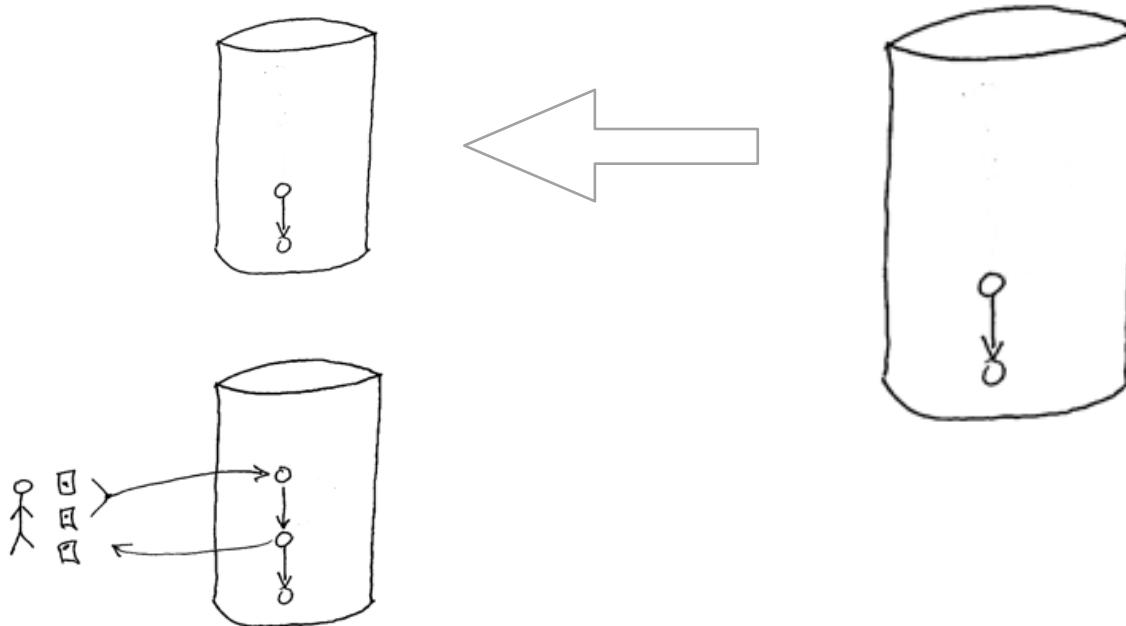
More

Merging



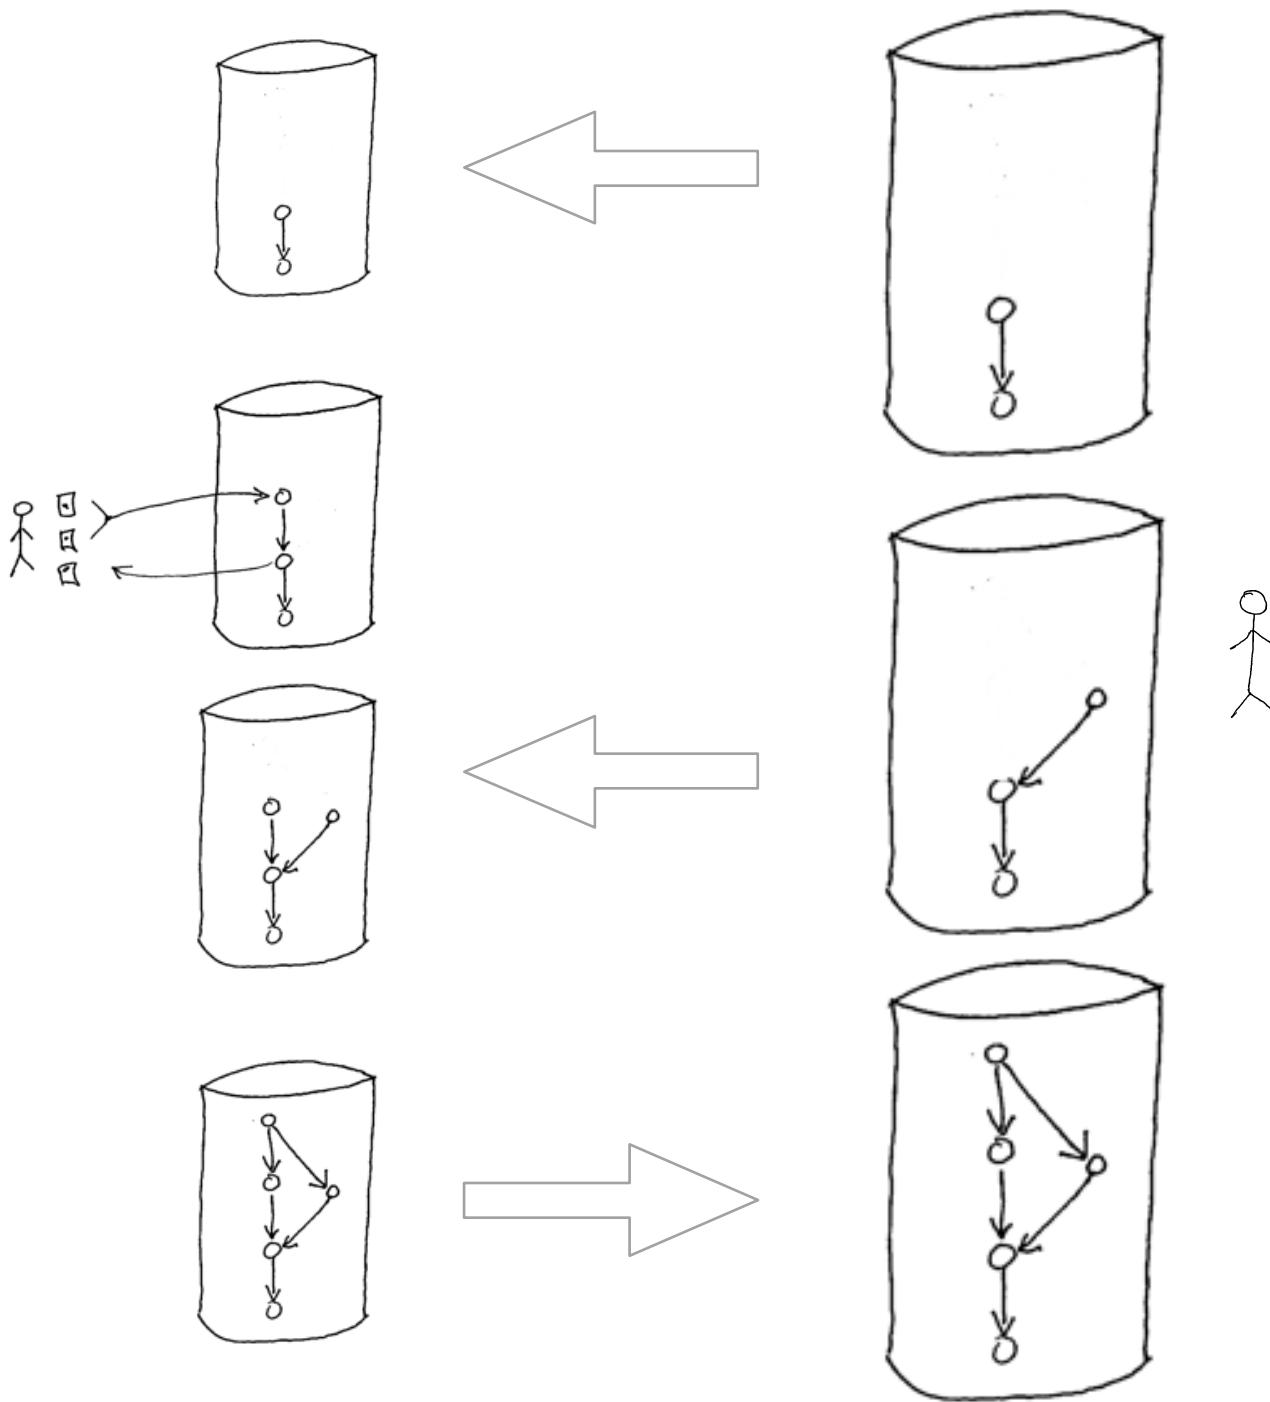
Author	Subject
Stuart.	added support for the freertos.armv7m target to the new_application script
Stuart.	created new_application script
Don G.	Merge pull request #1 from Railstars/master
Railstars	Merge branch 'master' of https://github.com/Railstars/openmrn
Stuart	added 64-bit MacOS target and fixed array overrun bug in grid c...
Stuart	reverted stellarisware makefile change
Stuart	Merge branch 'master' of https://github.com/Railstars/openmrn
Stuart	first attempt at a macos port
Railstars	Reverted an unnecessary change to the stellarisware makefile;...
Railstars	Added an include path for the Stellarisware library so that the Fre...
Railstars	Compiles, finally. Ran into an issue with recursive mutex's, think th...
Railstars	Changes to accomodate OS X's non-POSIX compliance re: clock_ge...
Stuart.	fixed typos and grammer mistakes
Stuart.	additional documentation and the addition of 3rd party XML and JSO...
Stuart.	bug fixes specific to the FreeRTOS/newlib port to get the dk-lm3s9...
Stuart.	added hook for board specific hardware initialization

Multiple repositories with easy transfer of commits between



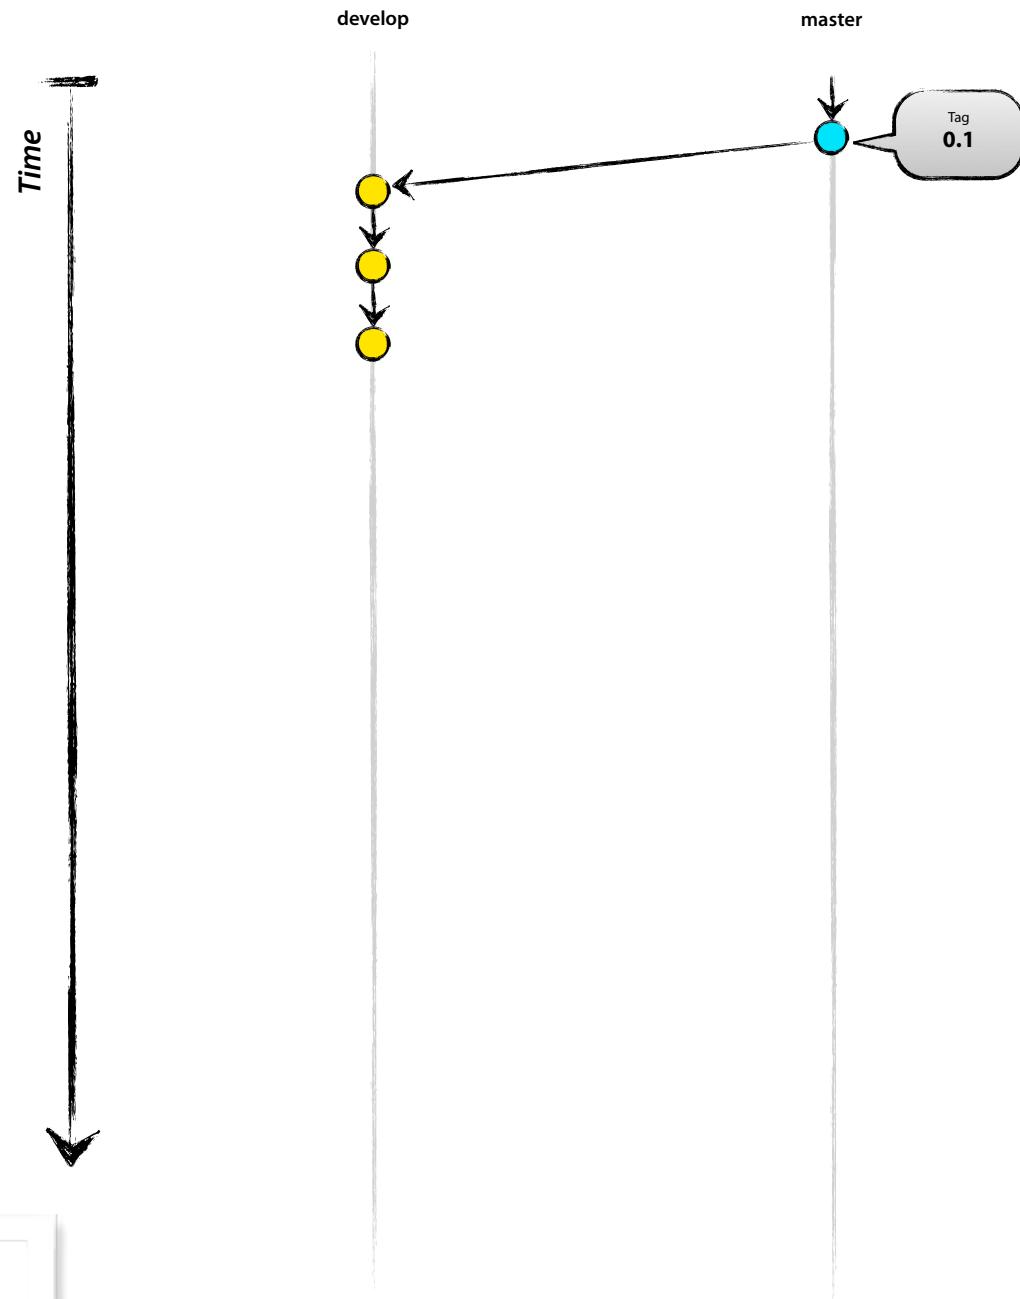
More

More than just mirroring



Branches are key

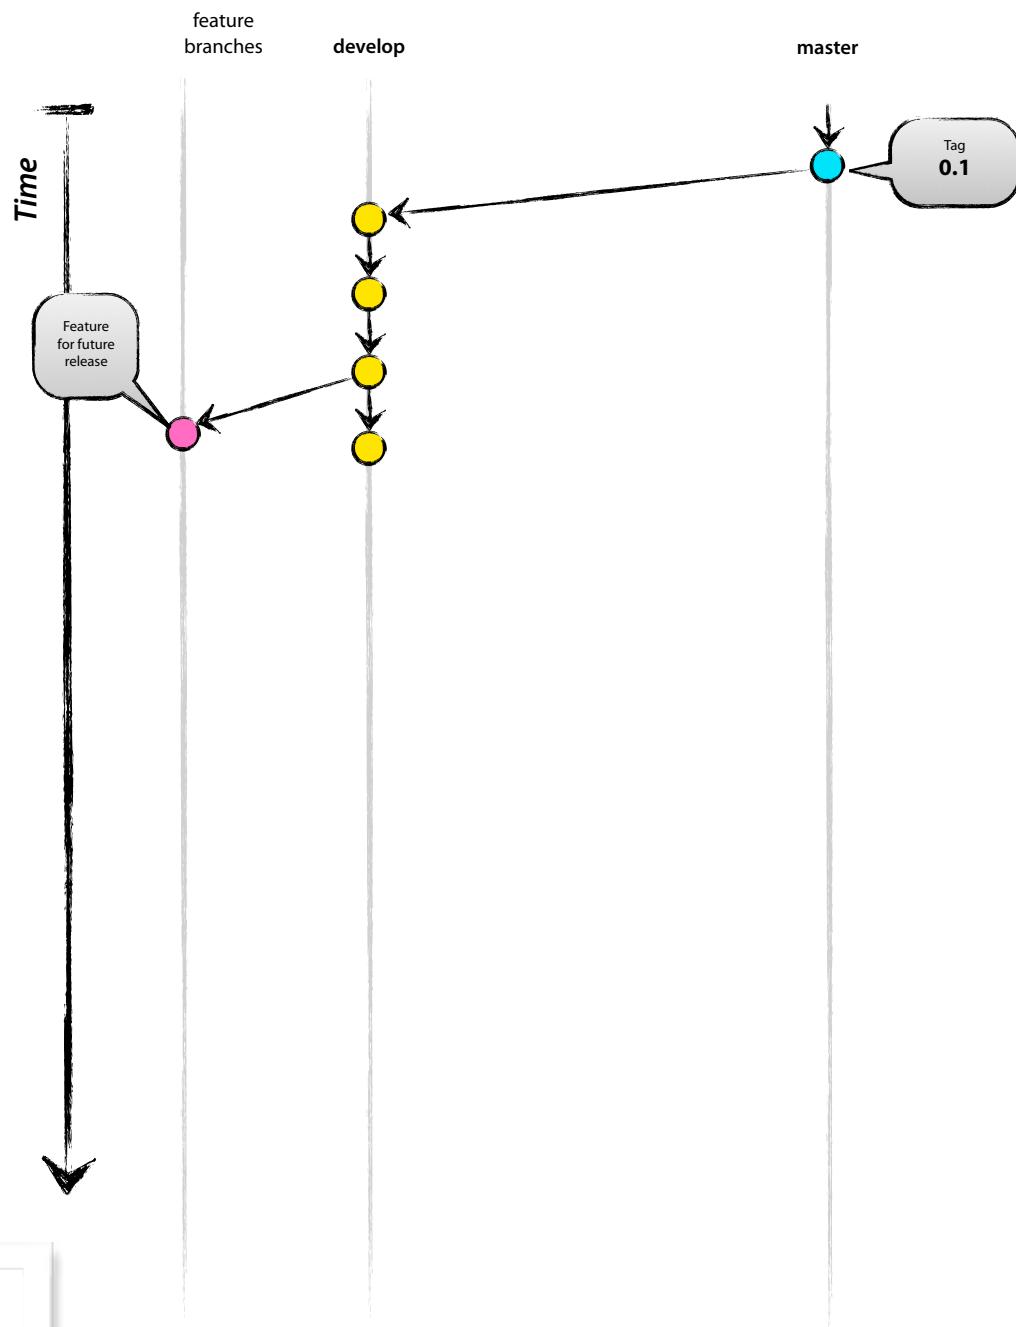
- Develop on a separate branch



Author: Vincent Driessen
Original blog post: <http://nvie.com/archives/323>
(cc) BY-SA

Branches are key

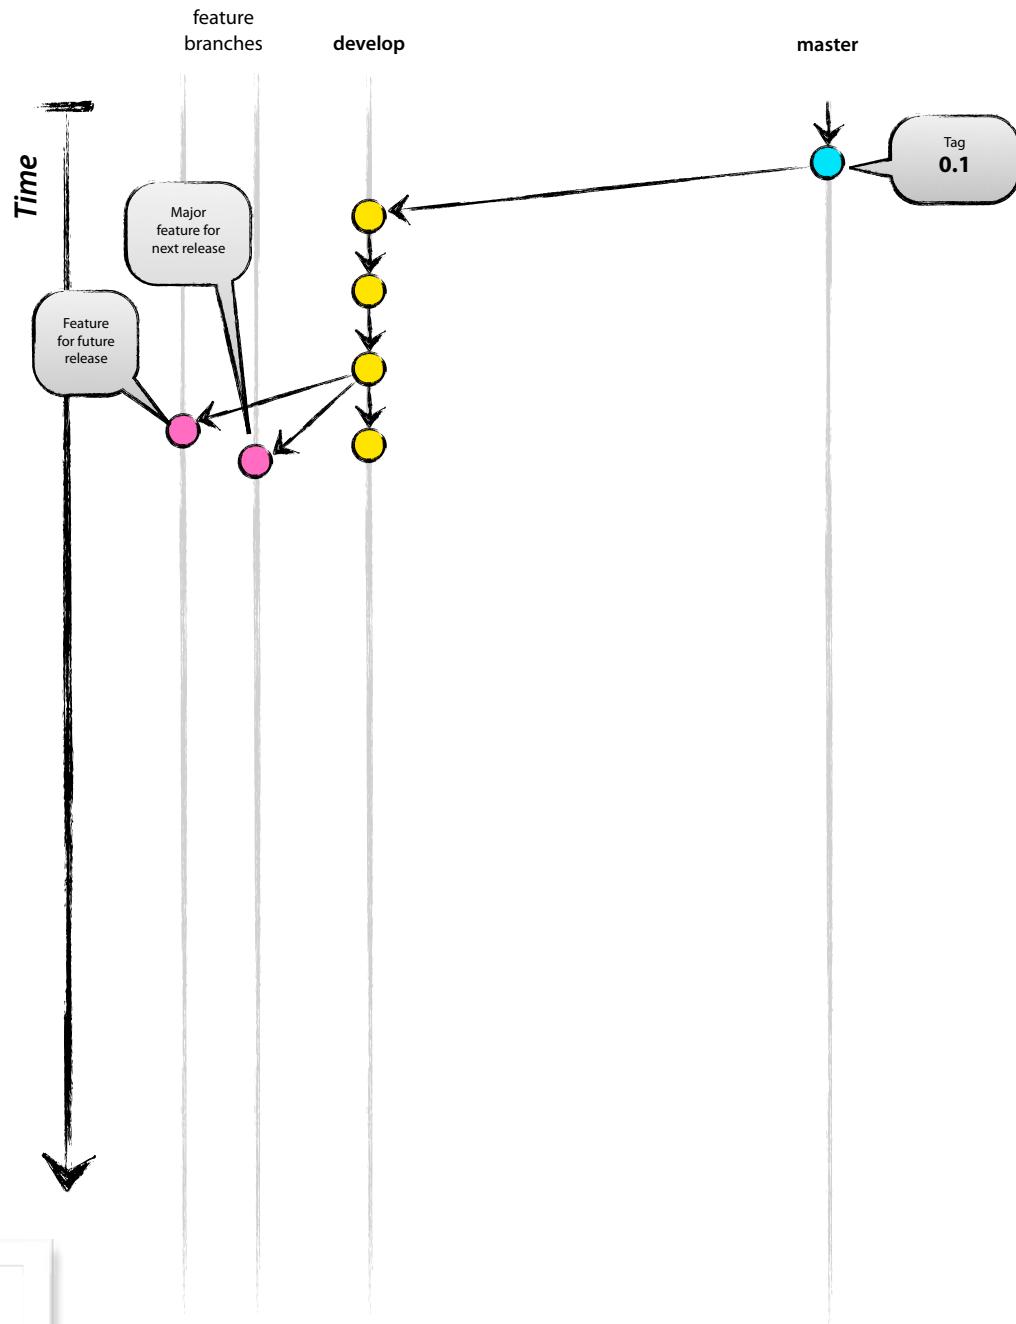
- Develop on a separate branch
- Future Big Feature on branch



Author: Vincent Driessens
Original blog post: <http://nvie.com/archives/323>
(cc) BY-SA

Branches are key

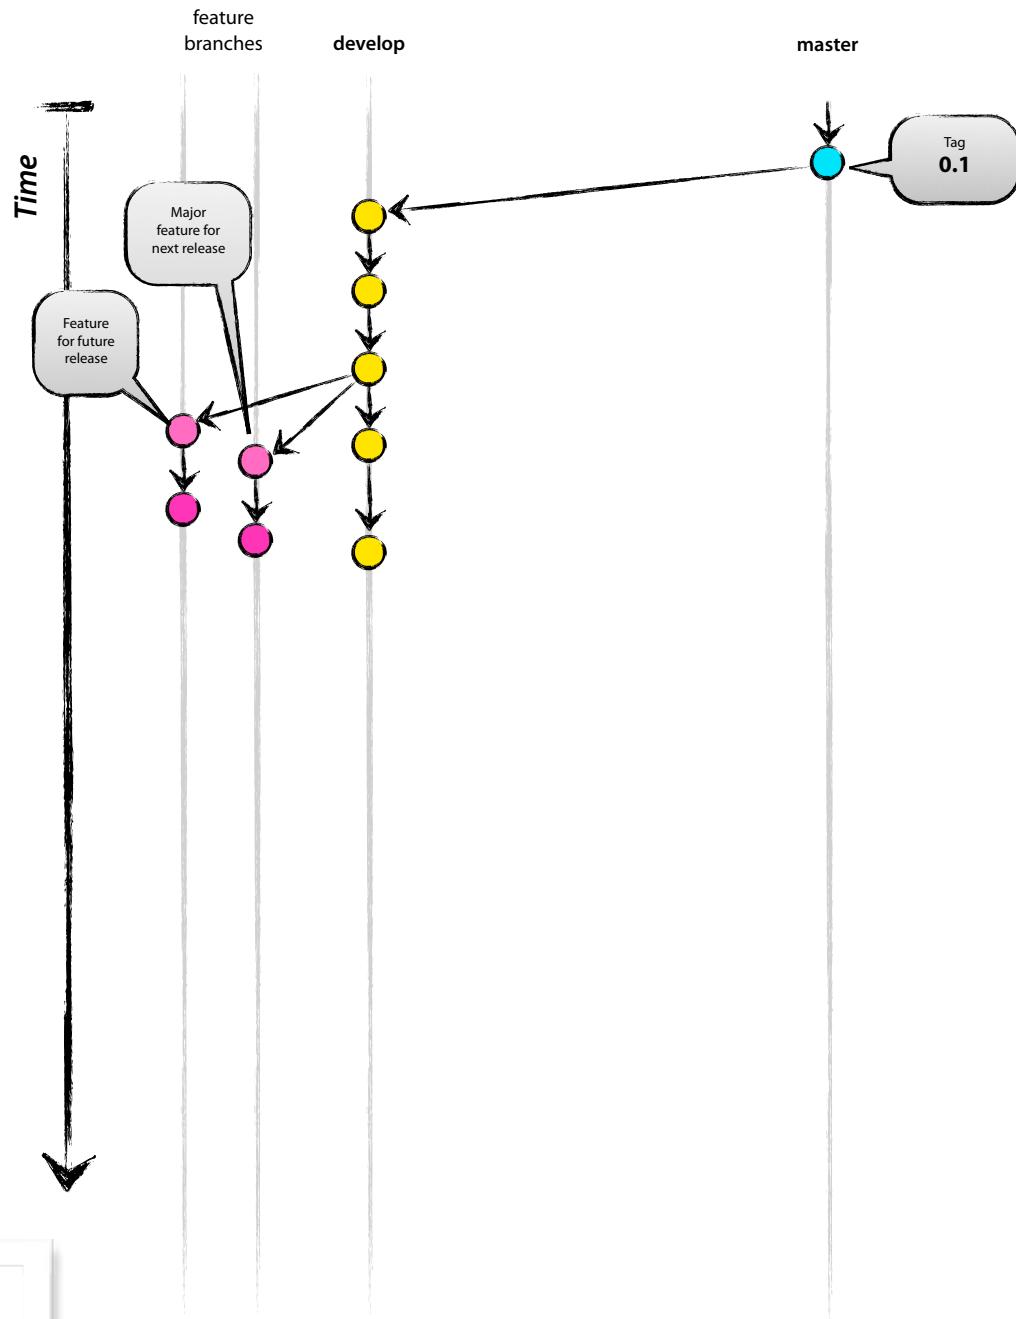
- Develop on a separate branch
- Future Big Feature on branch
- And another one



Author: Vincent Driessens
Original blog post: <http://nvie.com/archives/323>
(cc) BY-SA

Branches are key

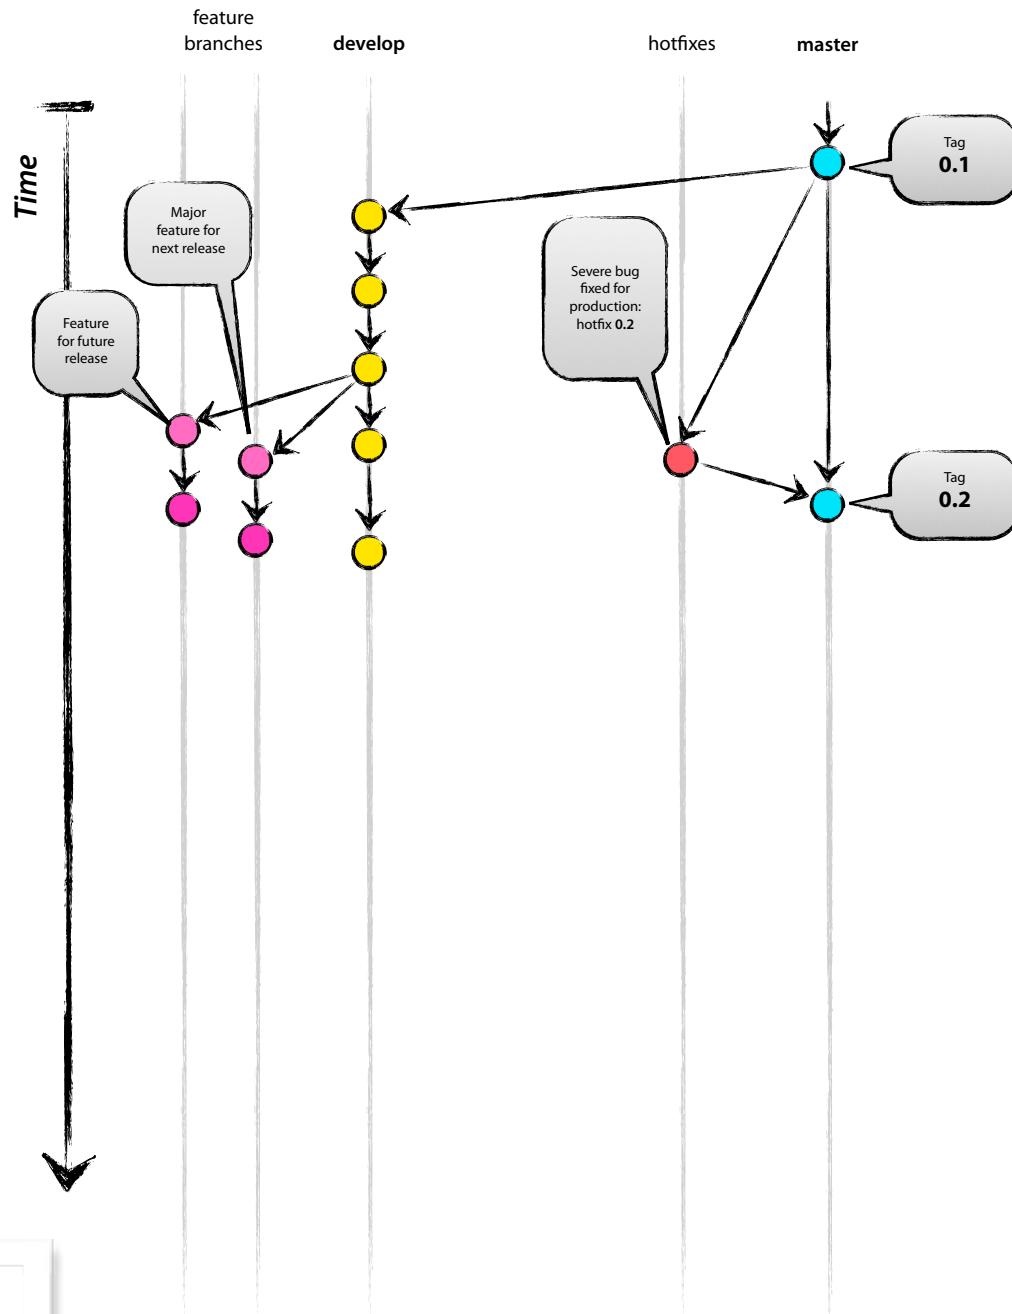
- Develop on a separate branch
- Future Big Feature on branch
- And another one for ll work



Author: Vincent Driessens
Original blog post: <http://nvie.com/archives/323>
(cc) BY-SA

Branches are key

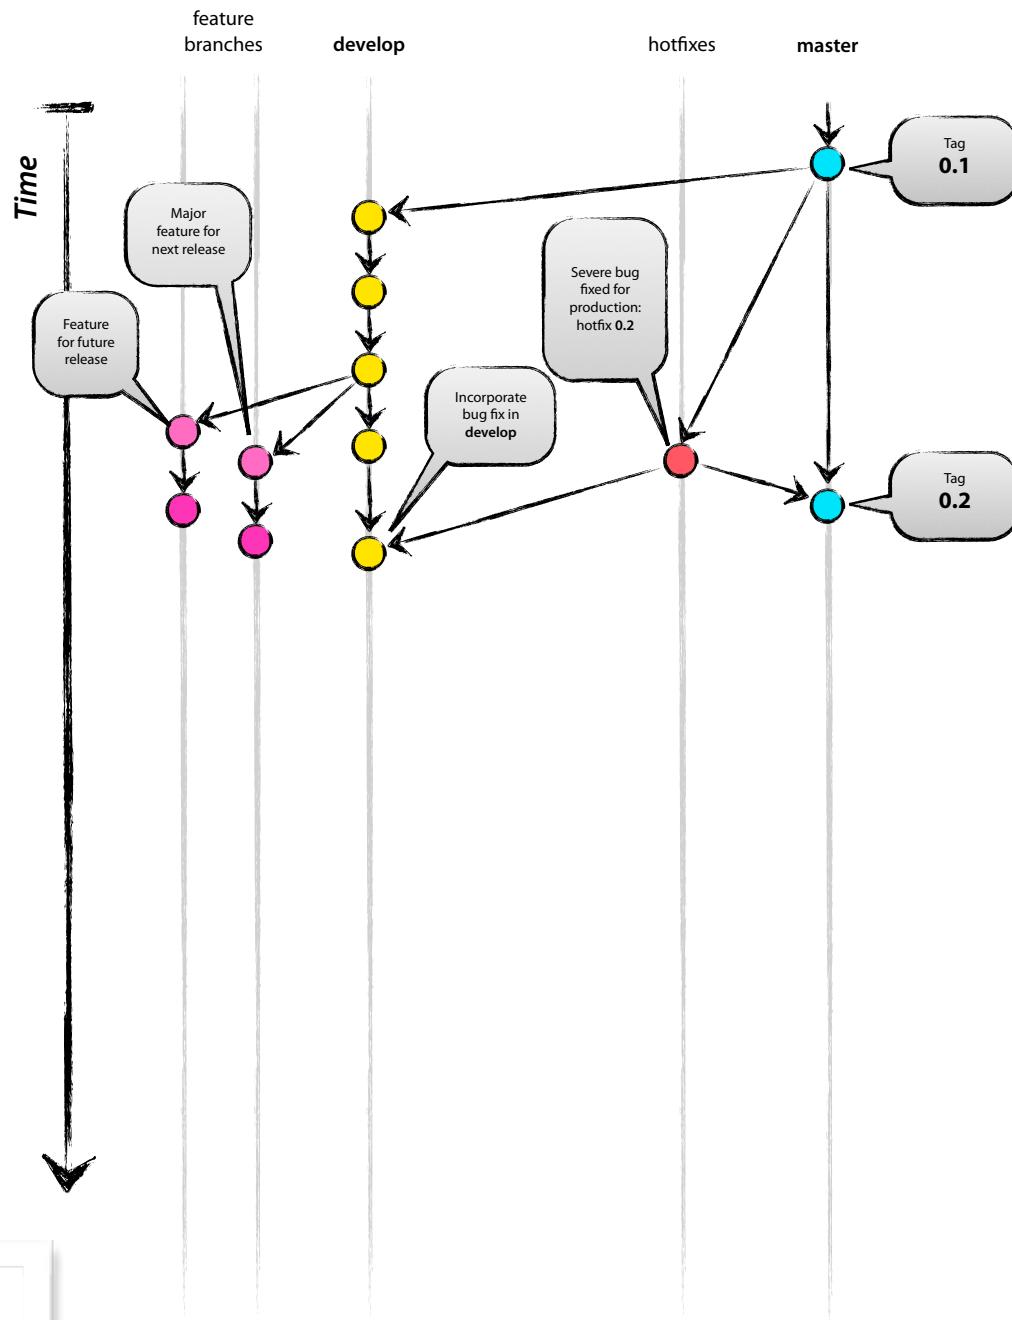
- Develop on a separate branch
- Future Big Feature on branch
- And another one for ll work
- Pays off for bug fix!



Author: Vincent Driessens
Original blog post: <http://nvie.com/archives/323>
(cc) BY-SA

Branches are key

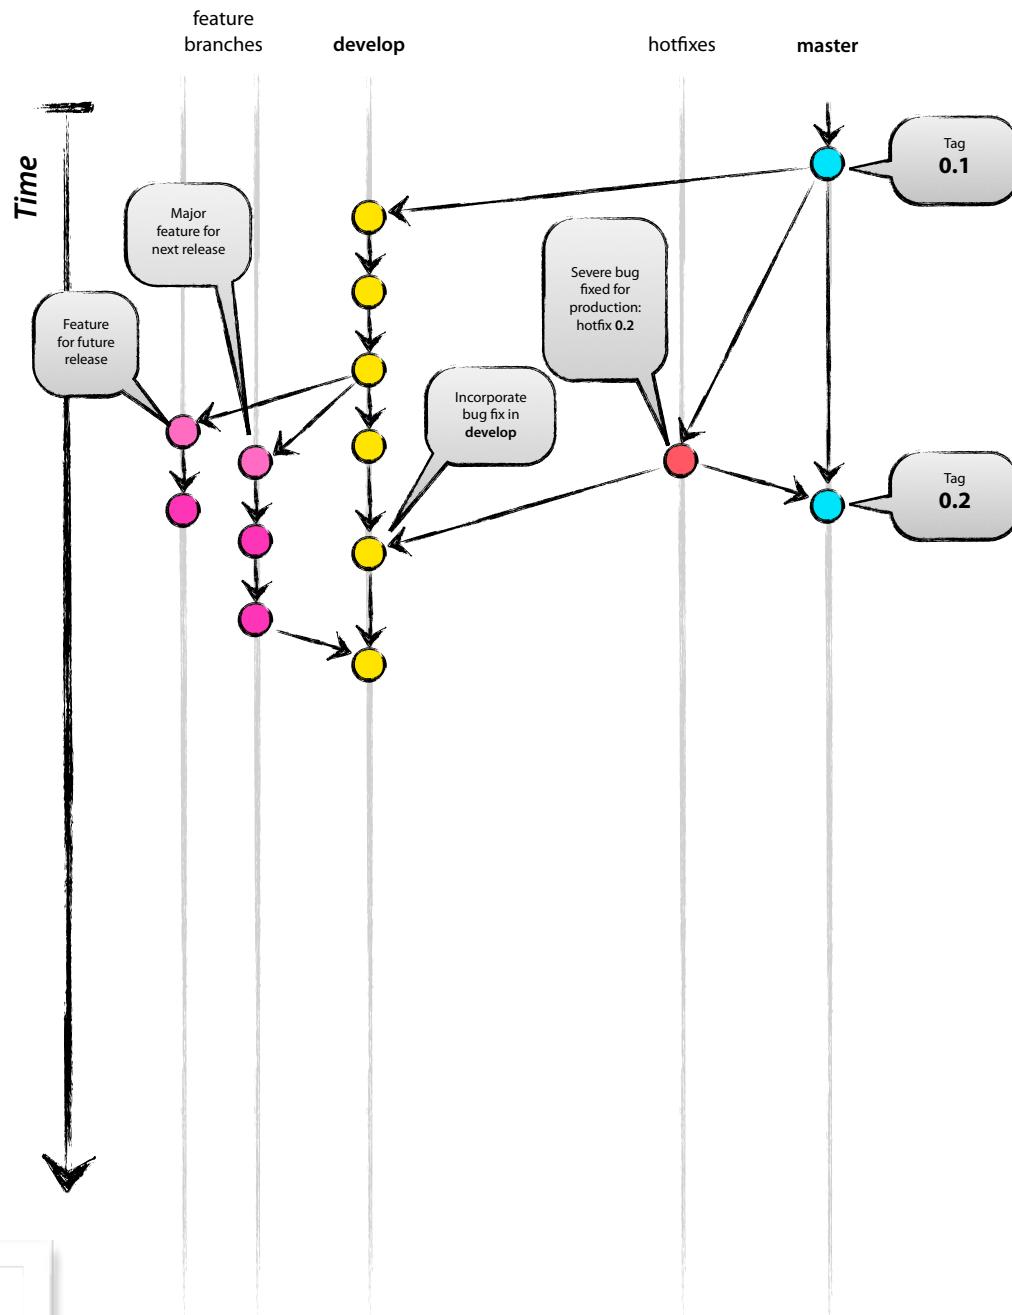
- Develop on a separate branch
- Future Big Feature on branch
- And another one for ll work
- Pays off for bug fix!
- Git merge to get fix across



Author: Vincent Driessens
 Original blog post: <http://nvie.com/archives/323>
 (cc) BY-SA

Branches are key

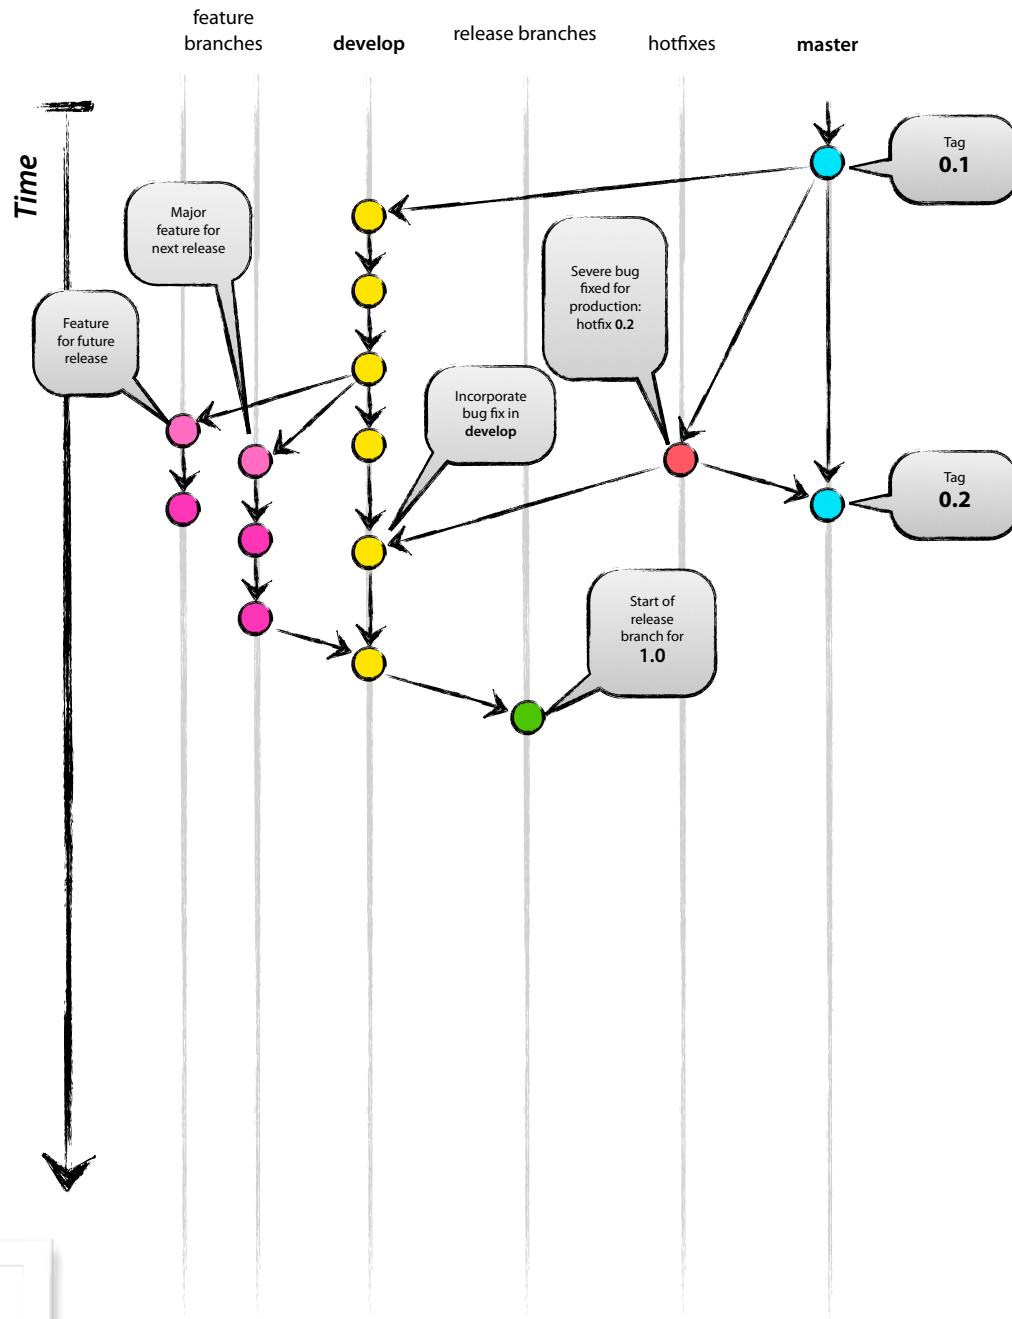
- Develop on a separate branch
- Future Big Feature on branch
- And another one for ll work
- Pays off for bug fix!
- Git merge to get fix across
- Feature done, merges in



Author: Vincent Driessens
 Original blog post: <http://nvie.com/archives/323>
 (cc) BY-SA

Branches are key

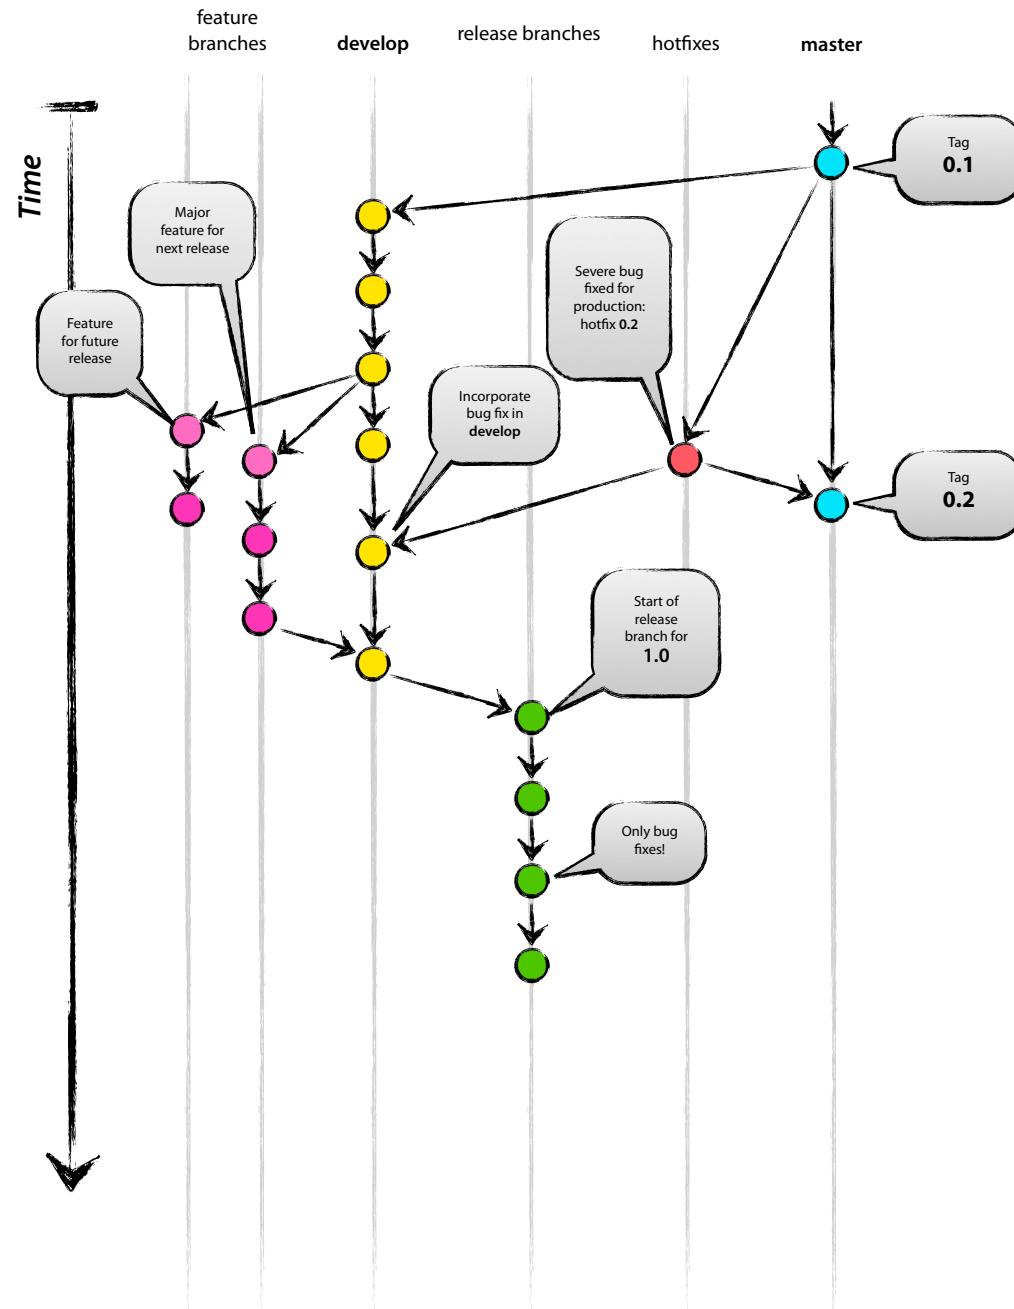
- Develop on a separate branch
- Future Big Feature on branch
- And another one for ll work
- Pays off for bug fix!
- Git merge to get fix across
- Feature done, merges in
- New branch holds release



Author: Vincent Driessens
 Original blog post: <http://nvie.com/archives/323>
 (cc) BY-SA

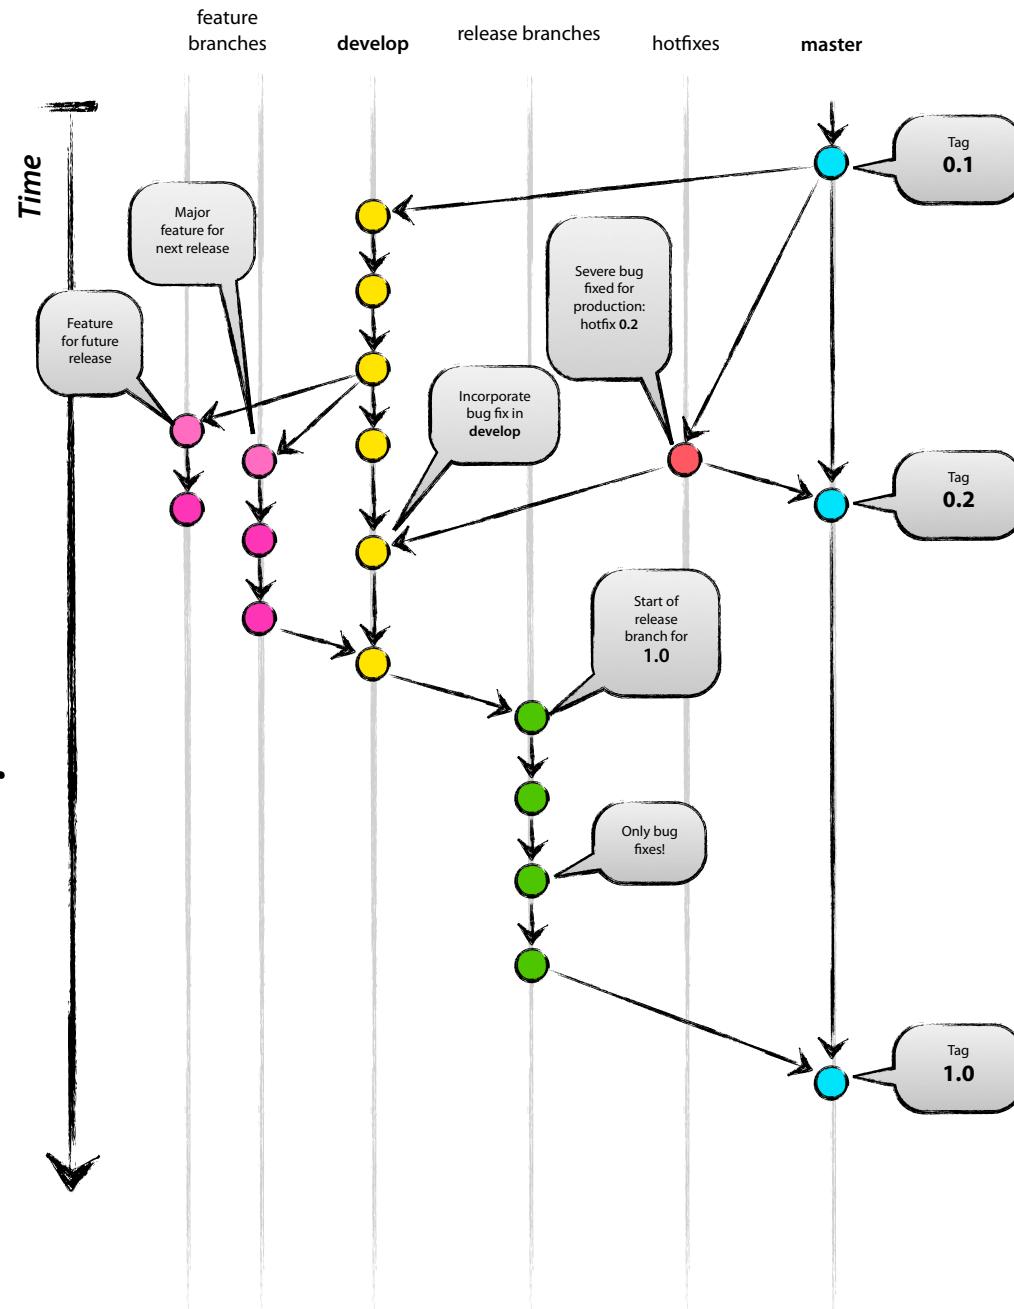
Branches are key

- Develop on a separate branch
- Future Big Feature on branch
- And another one for ll work
- Pays off for bug fix!
- Git merge to get fix across
- Feature done, merges in
- New branch holds release
- and it's inevitable fixes



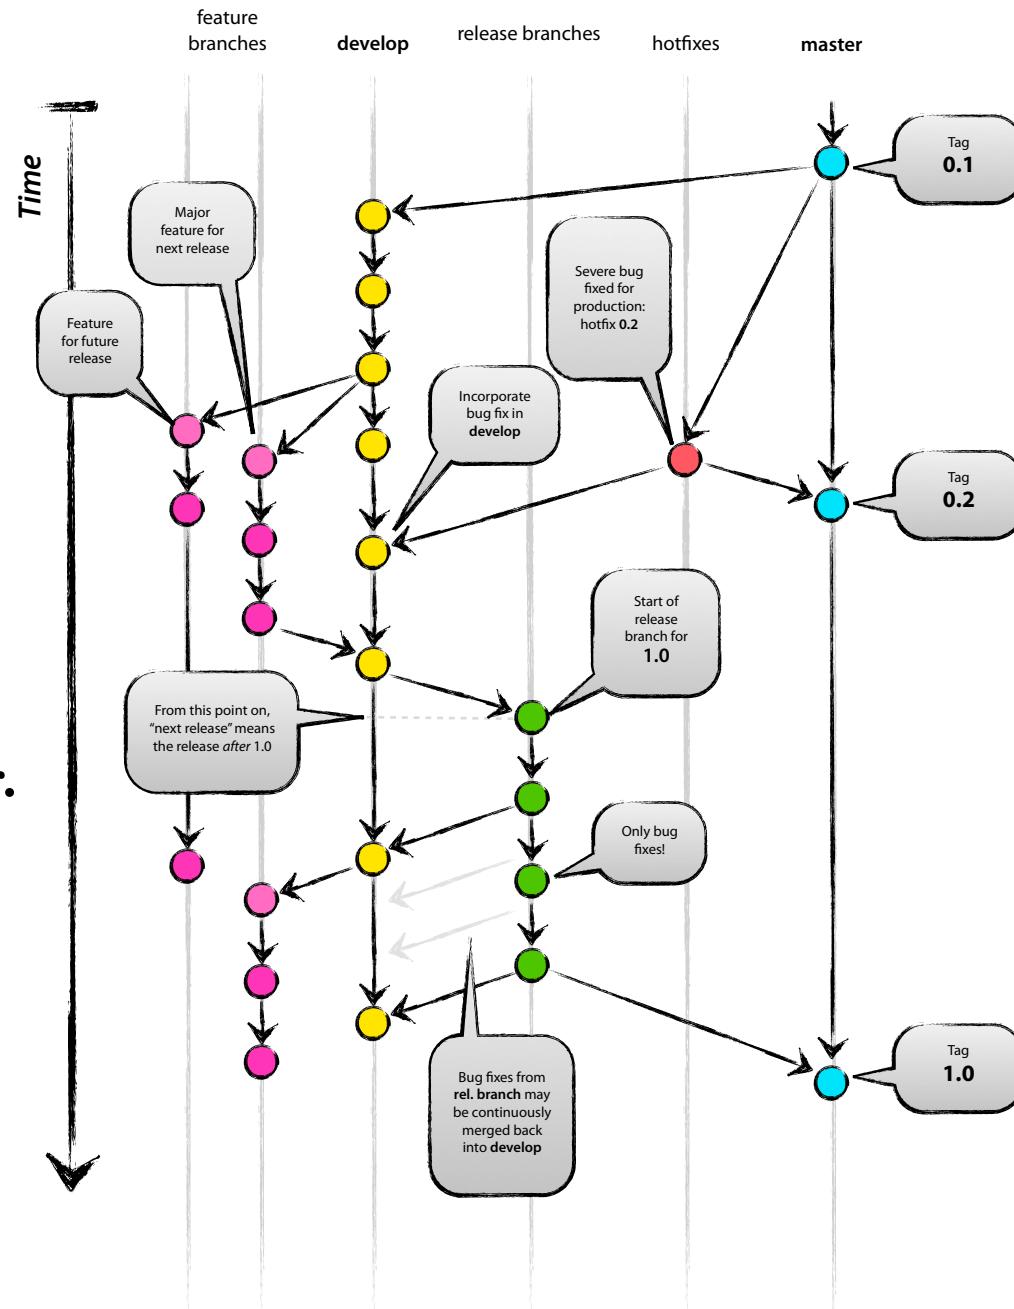
Branches are key

- Develop on a separate branch
- Future Big Feature on branch
- And another one for ll work
- Pays off for bug fix!
- Git merge to get fix across
- Feature done, merges in
- New branch holds release and it's inevitable fixes
- until merge and release master



Branches are key

- Develop on a separate branch
- Future Big Feature on branch
- And another one for ll work
- Pays off for bug fix!
- Git merge to get fix across
- Feature done, merges in
- New branch holds release and it's inevitable fixes until merge and release master.
- Meanwhile, work proceeds

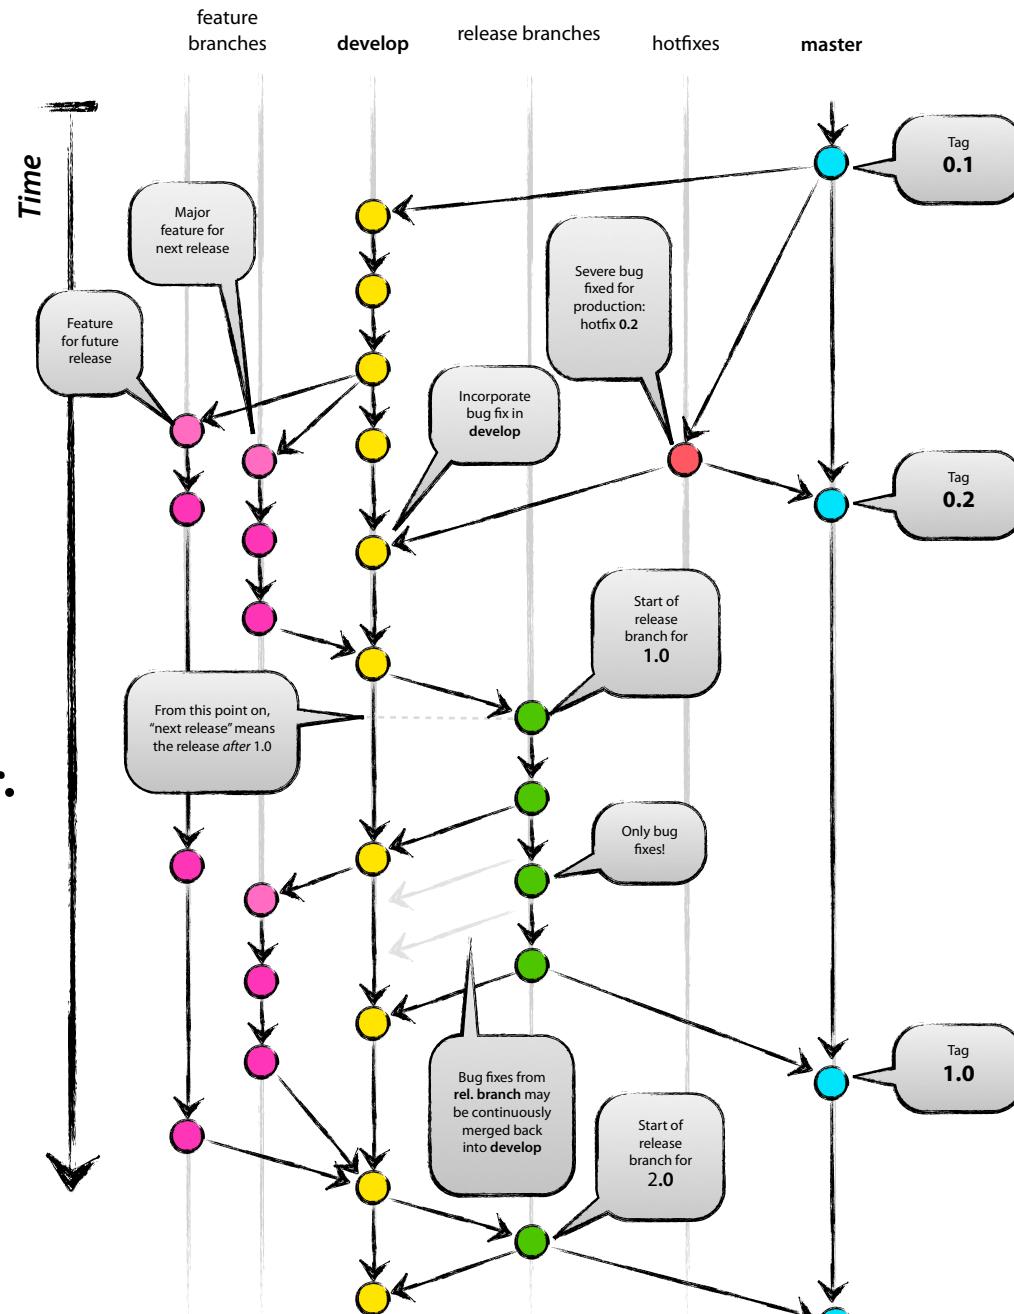


Branches are key

- Develop on a separate branch
- Future Big Feature on branch
- And another one for ll work
- Pays off for bug fix!
- Git merge to get fix across
- Feature done, merges in
- New branch holds release and it's inevitable fixes until merge and release master.
- Meanwhile, work proceeds
- And the process repeats

Keys: cheap branches,
reliable merges

Gives understandable story

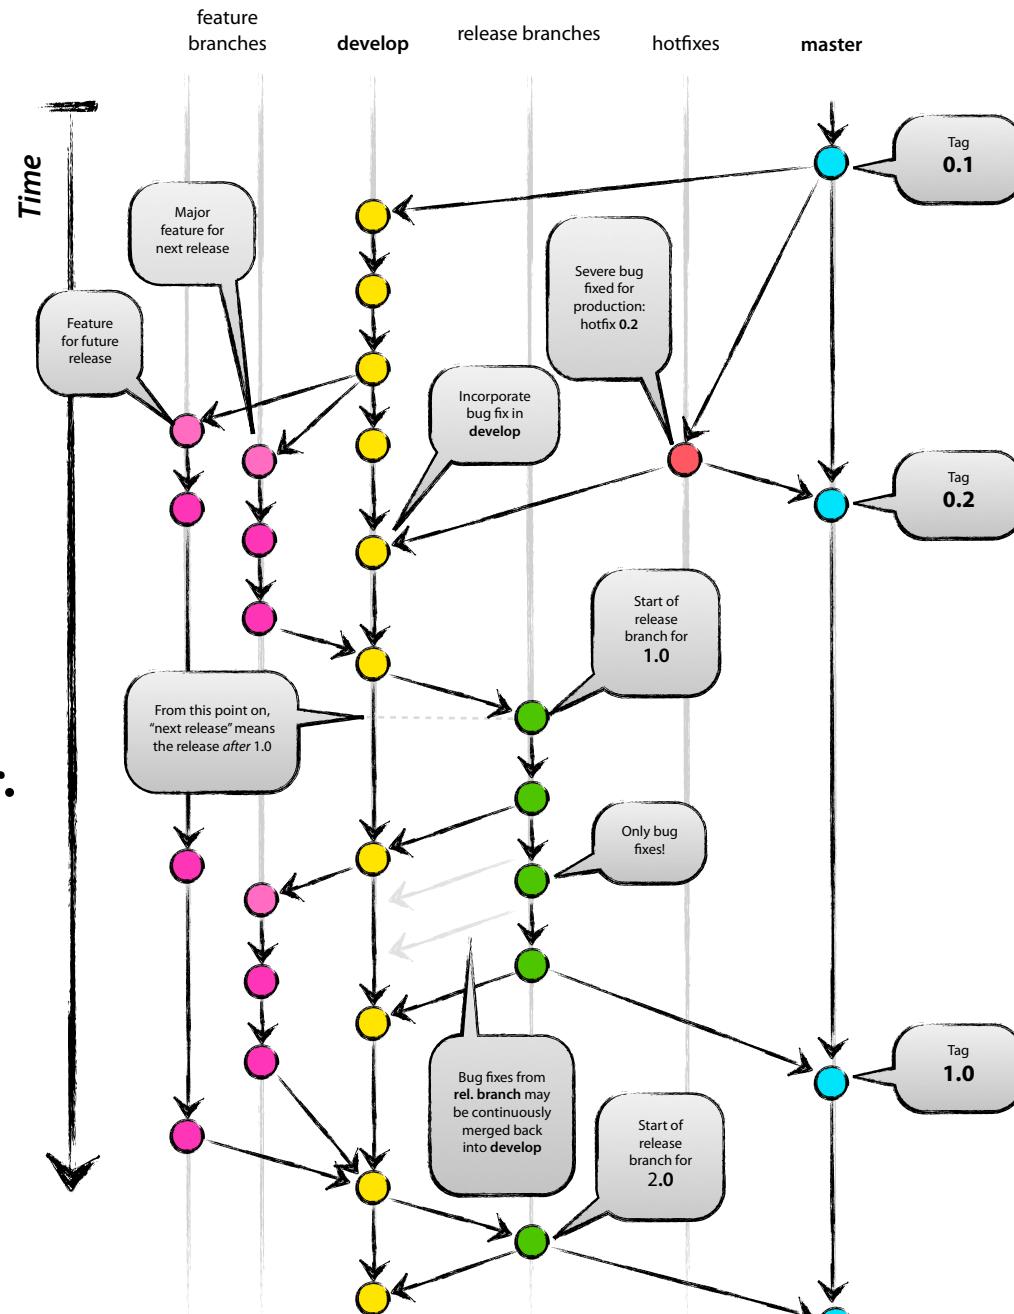


Branches are key

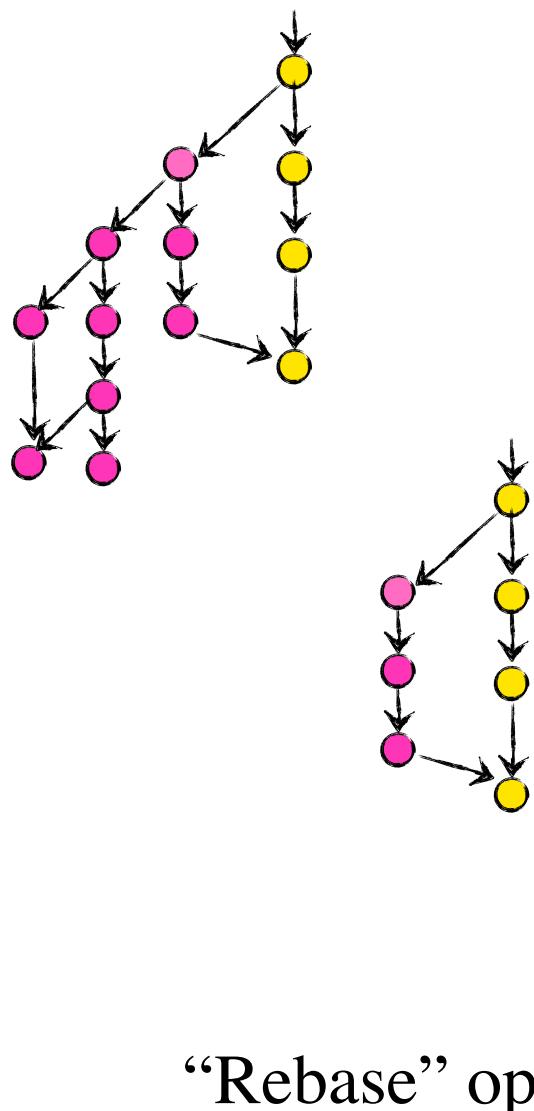
- Develop on a separate branch
- Future Big Feature on branch
- And another one for ll work
- Pays off for bug fix!
- Git merge to get fix across
- Feature done, merges in
- New branch holds release and it's inevitable fixes until merge and release master.
- Meanwhile, work proceeds
- And the process repeats

Keys: cheap branches,
reliable merges

Gives understandable story



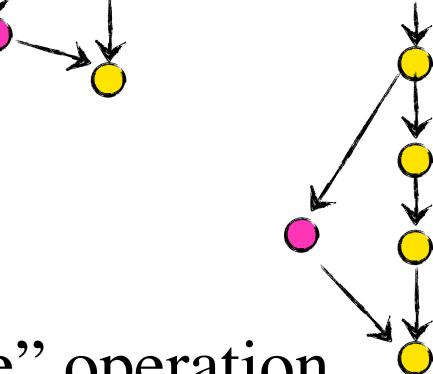
Rebase: An Editor for the Story



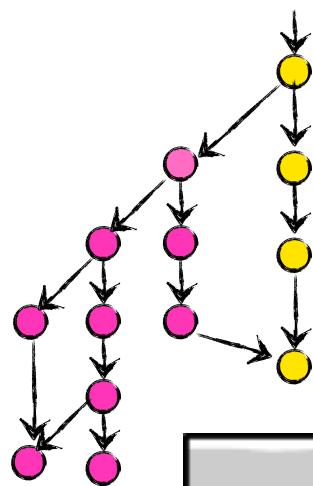
Finished difficult development task,
after several dead ends, lots of little
bits of progress & dead ends

More

Deleting only gets you so far



Rebase: An Editor for the Story



Finished difficult development task,
after several dead ends, lots of little
bits of progress & dead ends

More

	COMMENT	DATE
○	CREATED MAIN LOOP & TIMING CONTROL	14 HOURS AGO
○	ENABLED CONFIG FILE PARSING	9 HOURS AGO
○	MISC BUGFIXES	5 HOURS AGO
○	CODE ADDITIONS/EDITS	4 HOURS AGO
○	MORE CODE	4 HOURS AGO
○	HERE HAVE CODE	4 HOURS AGO
○	AAAAAAA	3 HOURS AGO
○	ADKFJSLKDFJSDFKLJ	3 HOURS AGO
○	MY HANDS ARE TYPING WORDS	2 HOURS AGO
○	HAAAAAAAAANDS	2 HOURS AGO

AS A PROJECT DRAGS ON, MY GIT COMMIT
MESSAGES GET LESS AND LESS INFORMATIVE.

You want me to trust how many people?

How do you give 6,000 people access to a central repository?

More

A) Don't! Have them submit patches to Package Coordinators (PBs)

```
Index: java/src/jmri/jmrix/openlcb/swing/downloader/LoaderPane.java
=====
--- java/src/jmri/jmrix/openlcb/swing/downloader/LoaderPane.java      (revision 29731)
+++ java/src/jmri/jmrix/openlcb/swing/downloader/LoaderPane.java      (working copy)
@@ -186,18 +186,19 @@
 */
void sendNext() {
    byte[] temp = new byte[SIZE];
-   int i;
-   for (i = 0; i < SIZE; i++) {
-       if (!inputContent.locationInUse(location+i))
+   int count;
+   for (count = 0; count < SIZE; count++) {
+       if (!inputContent.locationInUse(location+count)) {
+           break;
-       temp[i] = (byte)inputContent.getLocation(location+i);
+       }
+       temp[count] = (byte)inputContent.getLocation(location+count);
     }
-   byte[] data = new byte[i];
-   System.arraycopy(temp, 0, data, 0, i);
+   byte[] data = new byte[count];
System.arraycopy(temp, 0, data, 0, count);

    int addr = location; // next call back might be instantaneous
-   location = location + i;
-   log.info("Sending write to 0x{}", Integer.toHexString(location).toUpperCase());
+   location = location + count;
log.info("Sending write to 0x{} length {}", Integer.toHexString(location).toUpperCase(), count);
mcs.request(new MemoryConfigurationService.McsWriteMemo(destNodeID(), space, addr, data) {
    public void handleWriteReply(int code) {
        // update GUI intermittently

```



Enough info for reliable commit, but not a lot of context, and no reliable way to merge back if commit is delayed

You want me to trust how many people?

How do you give 6,000 people access to a central repository?

B) Find reliable people and give them access, log commits

```

Revision: 29733
http://sourceforge.net/p/jmri/code/29733
Author: jacobsen
Date: 2015-08-09 23:20:19 +0000 (Sun, 09 Aug 2015)
Log Message:
-----
Better index variable name; improve logging message

Modified Paths:
-----
trunk/jmri/java/src/jmri/jmrix/openlcb/swing/downloader/LoaderPane.java

Modified: trunk/jmri/java/src/jmri/jmrix/openlcb/swing/downloader/LoaderPane.java
=====
--- trunk/jmri/java/src/jmri/jmrix/openlcb/swing/downloader/LoaderPane.java      2015-08-08 23:10:01 UTC (rev 29732)
+++ trunk/jmri/java/src/jmri/jmrix/openlcb/swing/downloader/LoaderPane.java      2015-08-09 23:20:19 UTC (rev 29733)
@@ -186,18 +186,19 @@
 */
void sendNext() {
    byte[] temp = new byte[SIZE];
-   int i;
-   for (i = 0; i < SIZE; i++) {
-       if (linputContent.locationInUse(location+i))
+   int count;
+   for (count = 0; count < SIZE; count++) {
+       if (linputContent.locationInUse(location+count)) {
+           break;
-       temp[i] = (byte)inputContent.getLocation(location+i);
+       }
+       temp[count] = (byte)inputContent.getLocation(location+count);
    }
-   byte[] data = new byte[i];
-   System.arraycopy(temp, 0, data, 0, i);
+   byte[] data = new byte[count];
+   System.arraycopy(temp, 0, data, 0, count);

    int addr = location; // next call back might be instantaneous
-   location = location + i;
-   log.info("Sending write to 0x{}", Integer.toHexString(location).toUpperCase());
+   location = location + count;
+   log.info("Sending write to 0x{} length {}", Integer.toHexString(location).toUpperCase(), count);
    mcs.request(new MemoryConfigurationService.McsWriteMemo(destNodeID(), space, addr, data) {
        public void handleWriteReply(int code) {
            // update GUI intermittently

```



Solves to context & merge-back problem,
 but do you really have 6,000 reliable friends?

You want me to trust how many people?

How do you give 6,000 people access to a central repository?

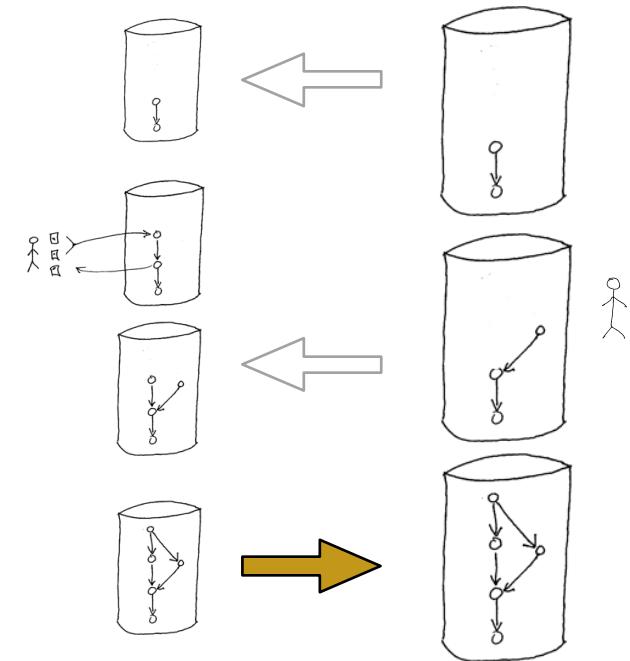
C) Use a distributed repository and “pull requests”

Git-based developers have a full local repository

Commits have full context

“Push” moves all that to target

A “pull request” sends all that to somebody at the target, who can accept or not



When accepted, the merge is completed & both repositories in sync

(Pull requests rarely rejected outright - usually it's “fix these things and resend”)

Strong tools exist to make pull requests easy: CI test results, etc automated

More

2nd approach: People handle consistency, machines build

Once consistency is managed, building is easier to automate

More

Enter “CMake”



Two phase process:

- (Zeroth: Pull a consistent set of code from managed repository)
- First, automatically build localized control files
- Second, do a platform-specific build using those files

```
cmake path  
make
```

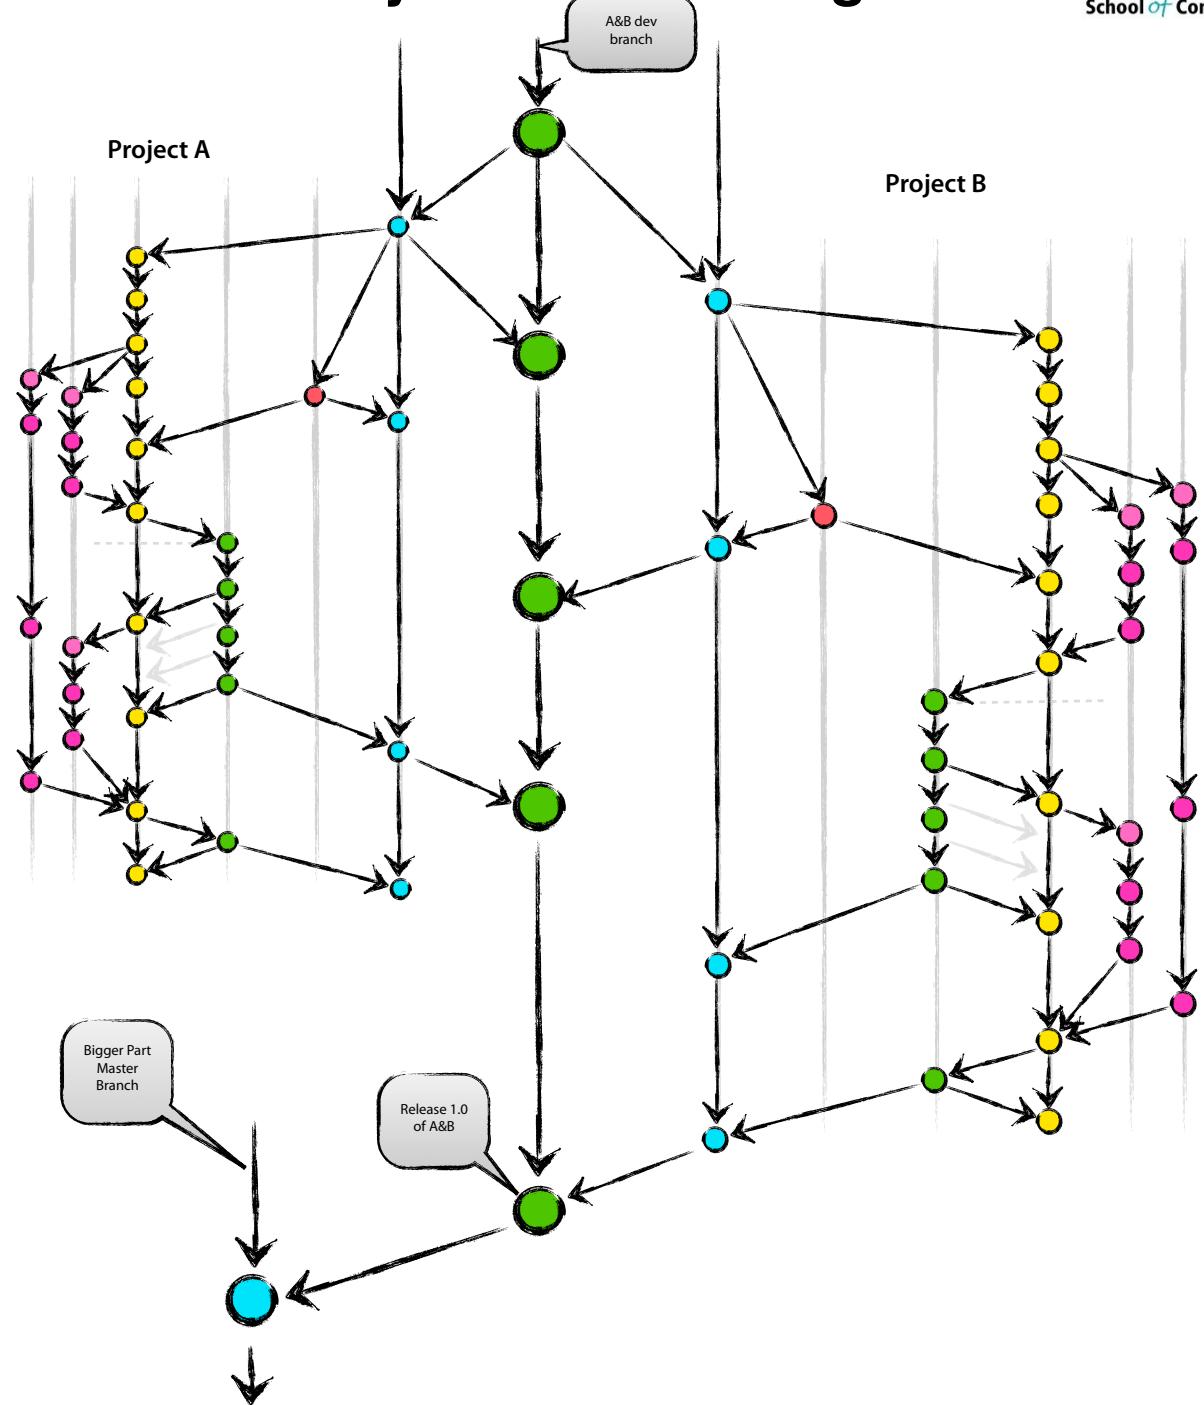
CMake:

- Scales very well (builds entire Linux distributions, LCGsoft LHC software)
- Well integrated with other tools (Eclipse, Visual Studio, the whole world)
- Powerful capabilities

Are you sure I don't need a version system for packages?

We use version control outside SVN because it's too hard to have lots of independent, controlled versions inside SVN.

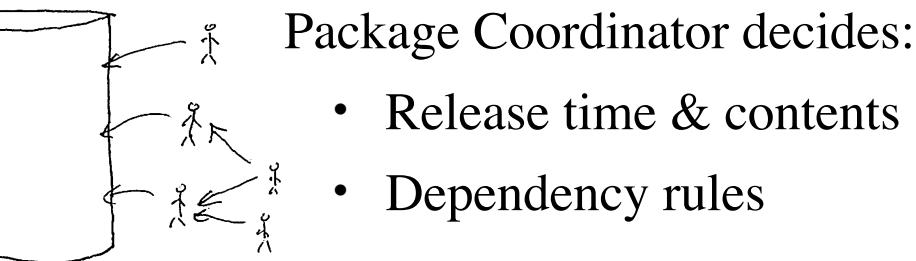
Git's “Lots of branches”
+
strong & easy merging
is
qualitatively different



Compare Approaches:

SVN & CMT

- **Code in repository**
- **Unit of organization: Package**



- **Tools create releases**

Resolve dependencies

- At package level

Specify localization

- **Build and distribute**

Pre-made Makefiles

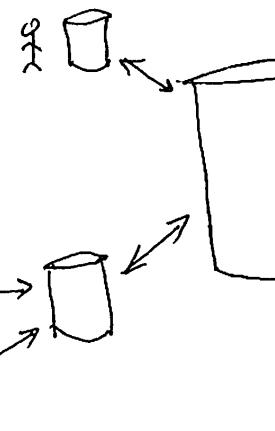
Git & CMake

- **Code in repository**
- **Unit of organization: Branch**

Fractal organization within

- Common time, contents
- Dependency implicit

Cooperate to define release



- **Check out and build**

Consistent release from branch

CMake handles localization

Exercises

Test Frameworks

Performance Profiling

Memory Issues

Code Management

Release Management

Instructions at: <http://server/jake/CSC/index.html>

More

You'll work in pairs. Try to find somebody with complementary skills!

Learn about each topic, spend more time on the ones that interest you.

Speed is not the issue: no reward for first done, no complaint about last.

Think about what you're doing: There are larger lessons to be found!

Lecture summary

Software engineering is the art of building complex computer systems

It's ideas and techniques spring from our need to handle size & complexity

As you do your own work & develop your own skills, consider:

- How your effort effects or contributes to things 10X, 100X, 1000X larger
- How you'll do things different/better when it's your problem

